

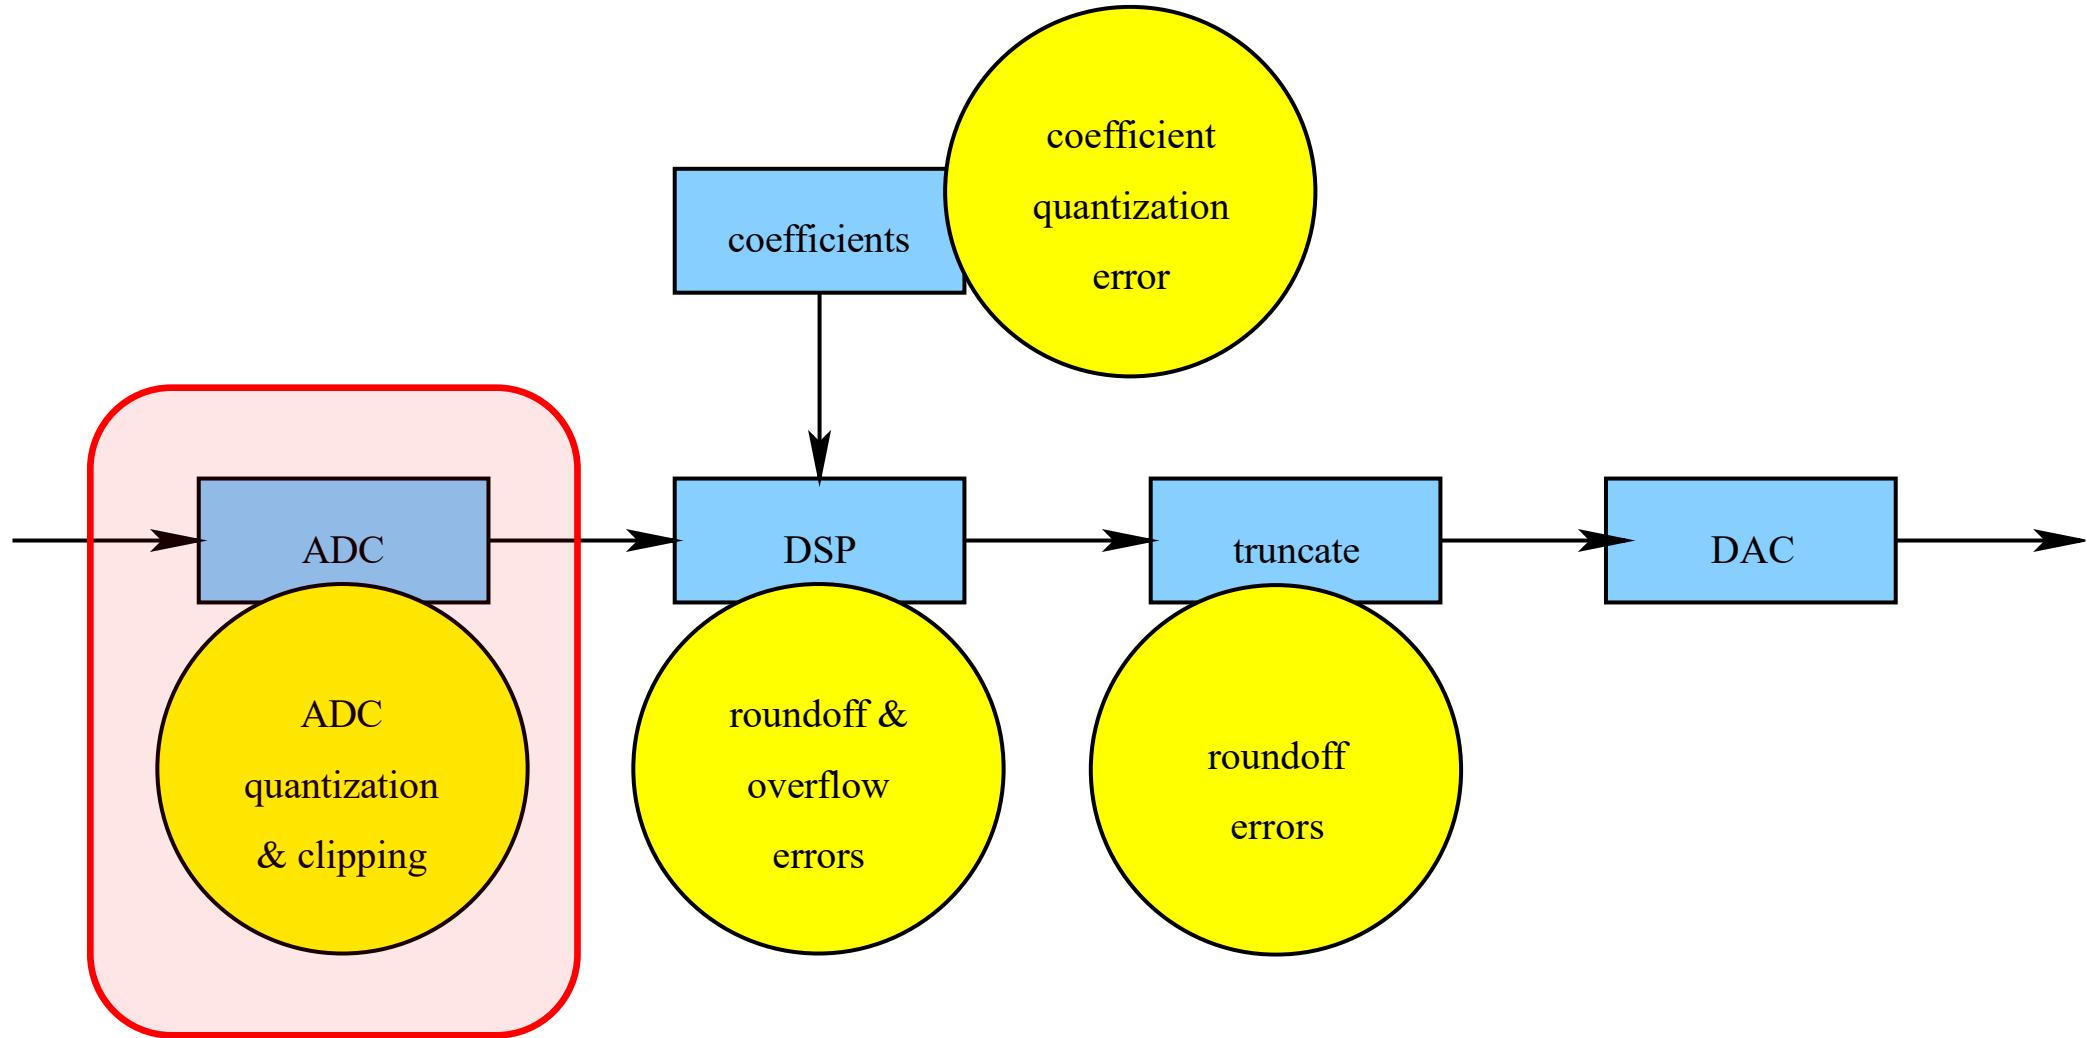
A/D conversion as an error source. Fixed-point IIR filter design.

Signal Processing Systems Fall 2025
Lecture 5 (Monday 10.11.)

Outline

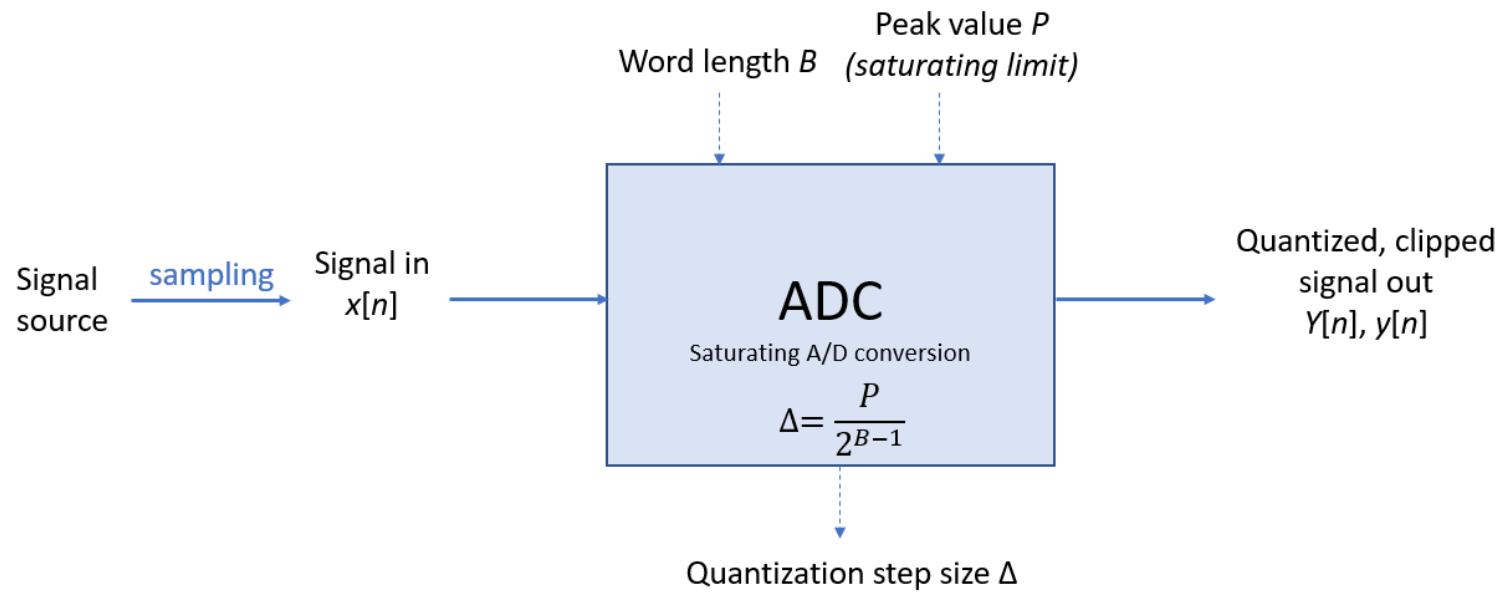
- A/D conversion as a noise source
 - Modelling ADC
 - Computing powers (using models / samples)
 - Crest factor, input noise floor, noise spectrum
- Fixed-point IIR filter design
 - Coefficient quantization & filter structure
 - Scoping of numerical ranges using Matlab's Fixed-Point Toolset
 - Second-order section (SOS) design
 - Input scaling
 - Experiments (with relation to ADC noise floor)

1. A/D conversion as a noise source



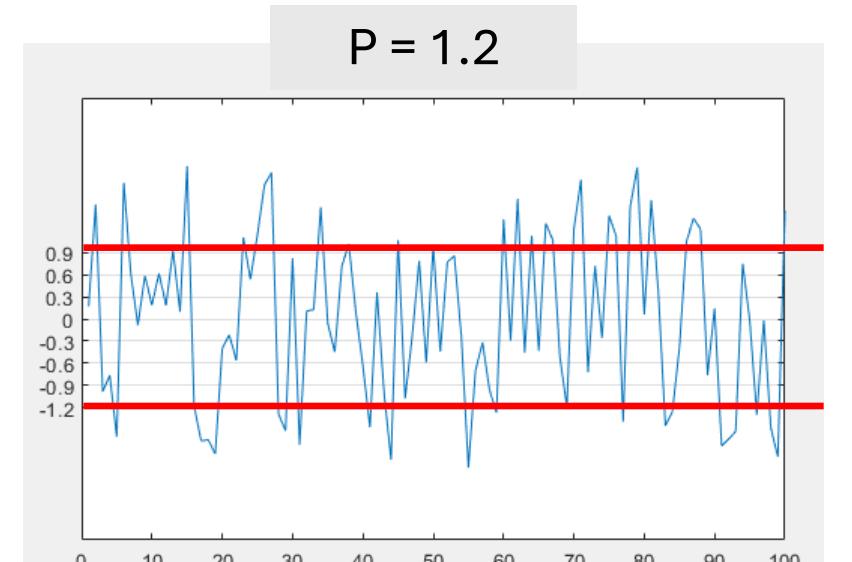
1.1. Model of A/D conversion

Saturating ADC:

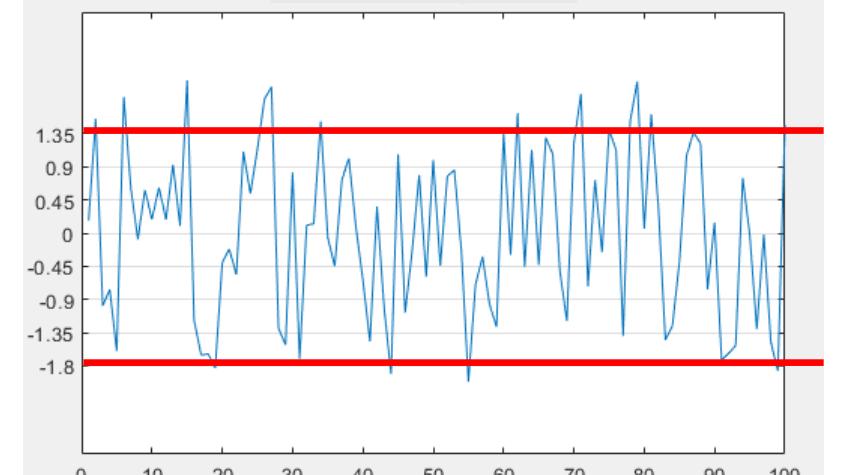


Larger $P \Rightarrow$ less clipping, but larger Δ

$B=3$ bits



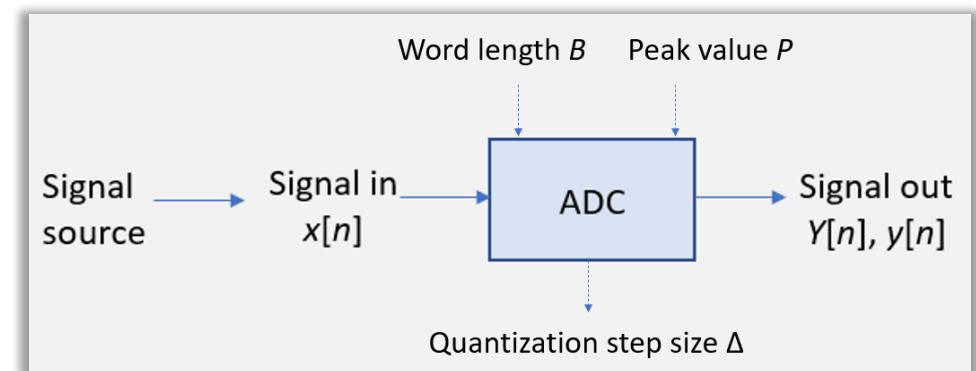
$\Delta = 0.3$ V



$\Delta = 0.45$ V

Model components

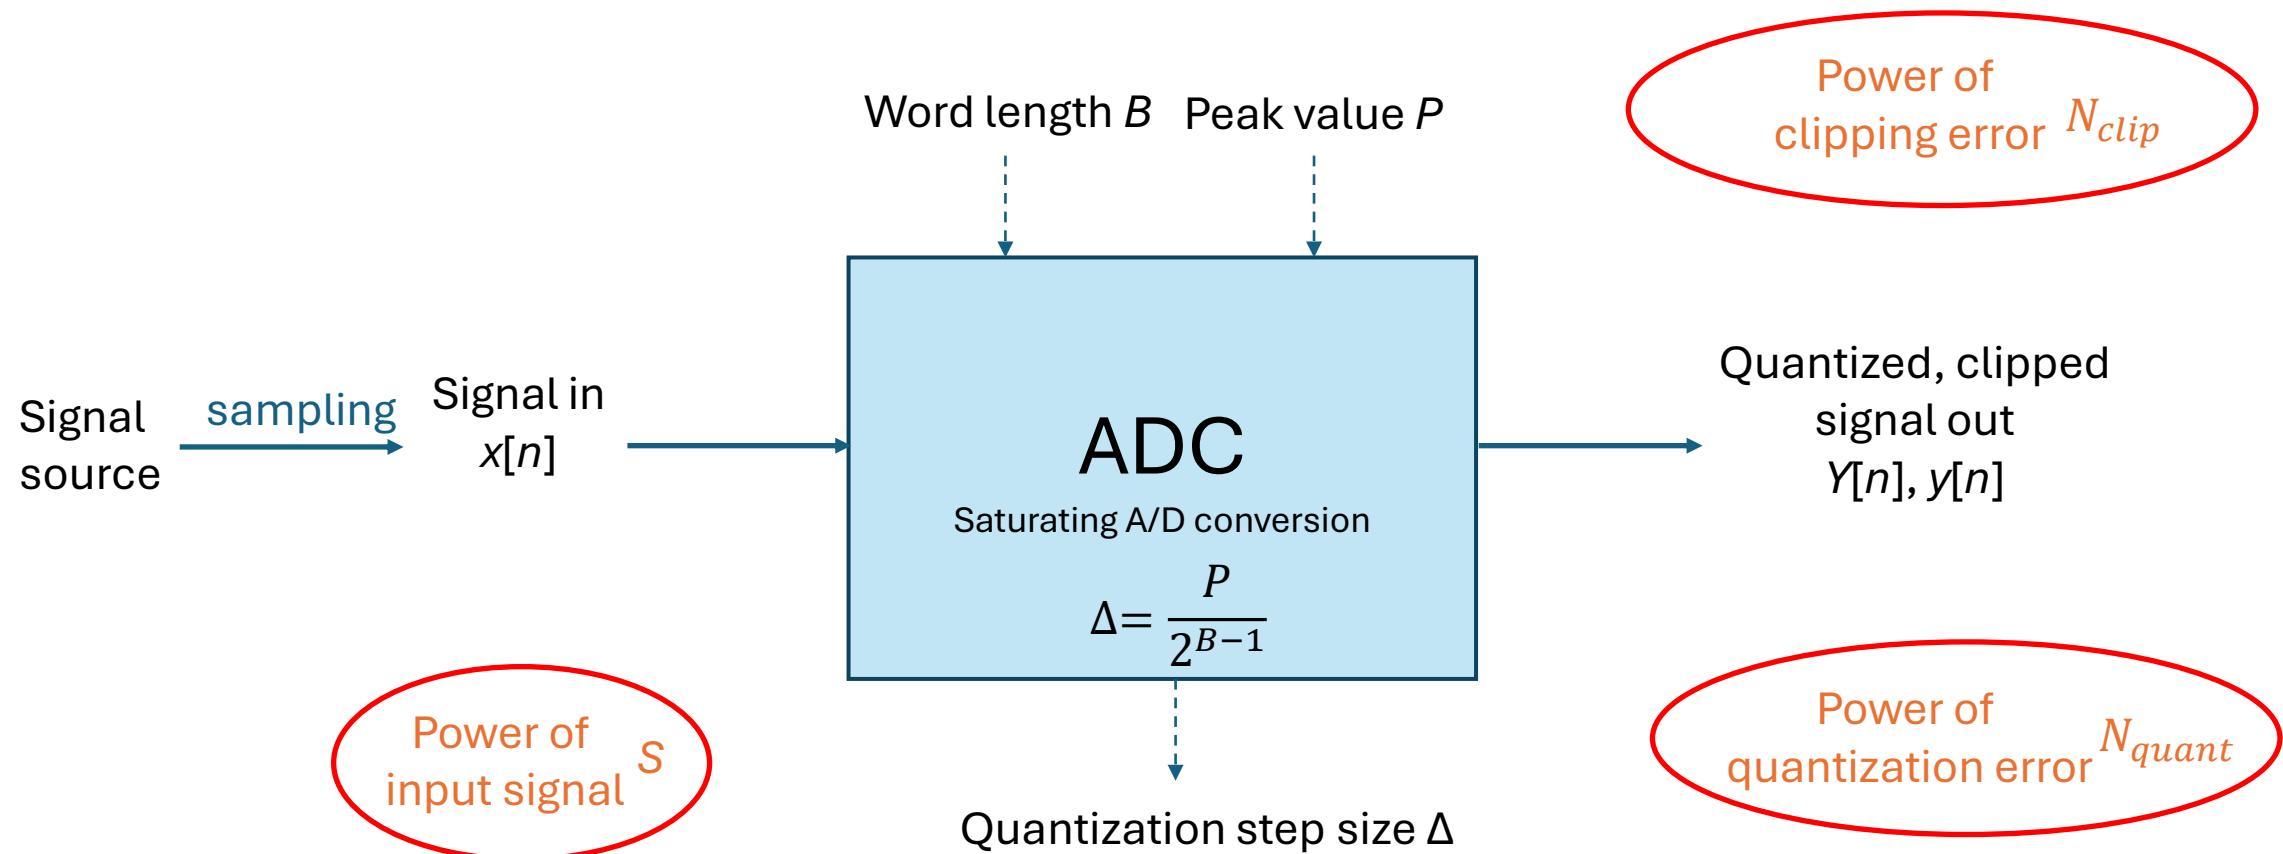
- Signal source
 - has certain characteristics, some quantities like power of the signal may be known (e.g. consider sine wave or Gaussian input)
- Input signal
 - sequence of samples, $x[n]$
- Output signal (quantized, clipped)
 - represented as two's complements integers $Y[n]$ (B bits)
 - values can be interpreted in terms of input signal levels, $y[n] = \Delta Y[n]$, where Δ is the quantization step size
- Saturating ADC
 - peak absolute value P of input signal that clipping does not occur
 - quantization step size Δ can be calculated from P and B



$$\Delta = \frac{P}{2^{B-1}}$$

Powers of interest

Goal: understand how signal/noise ratios change when B and P are adjusted



Computing powers

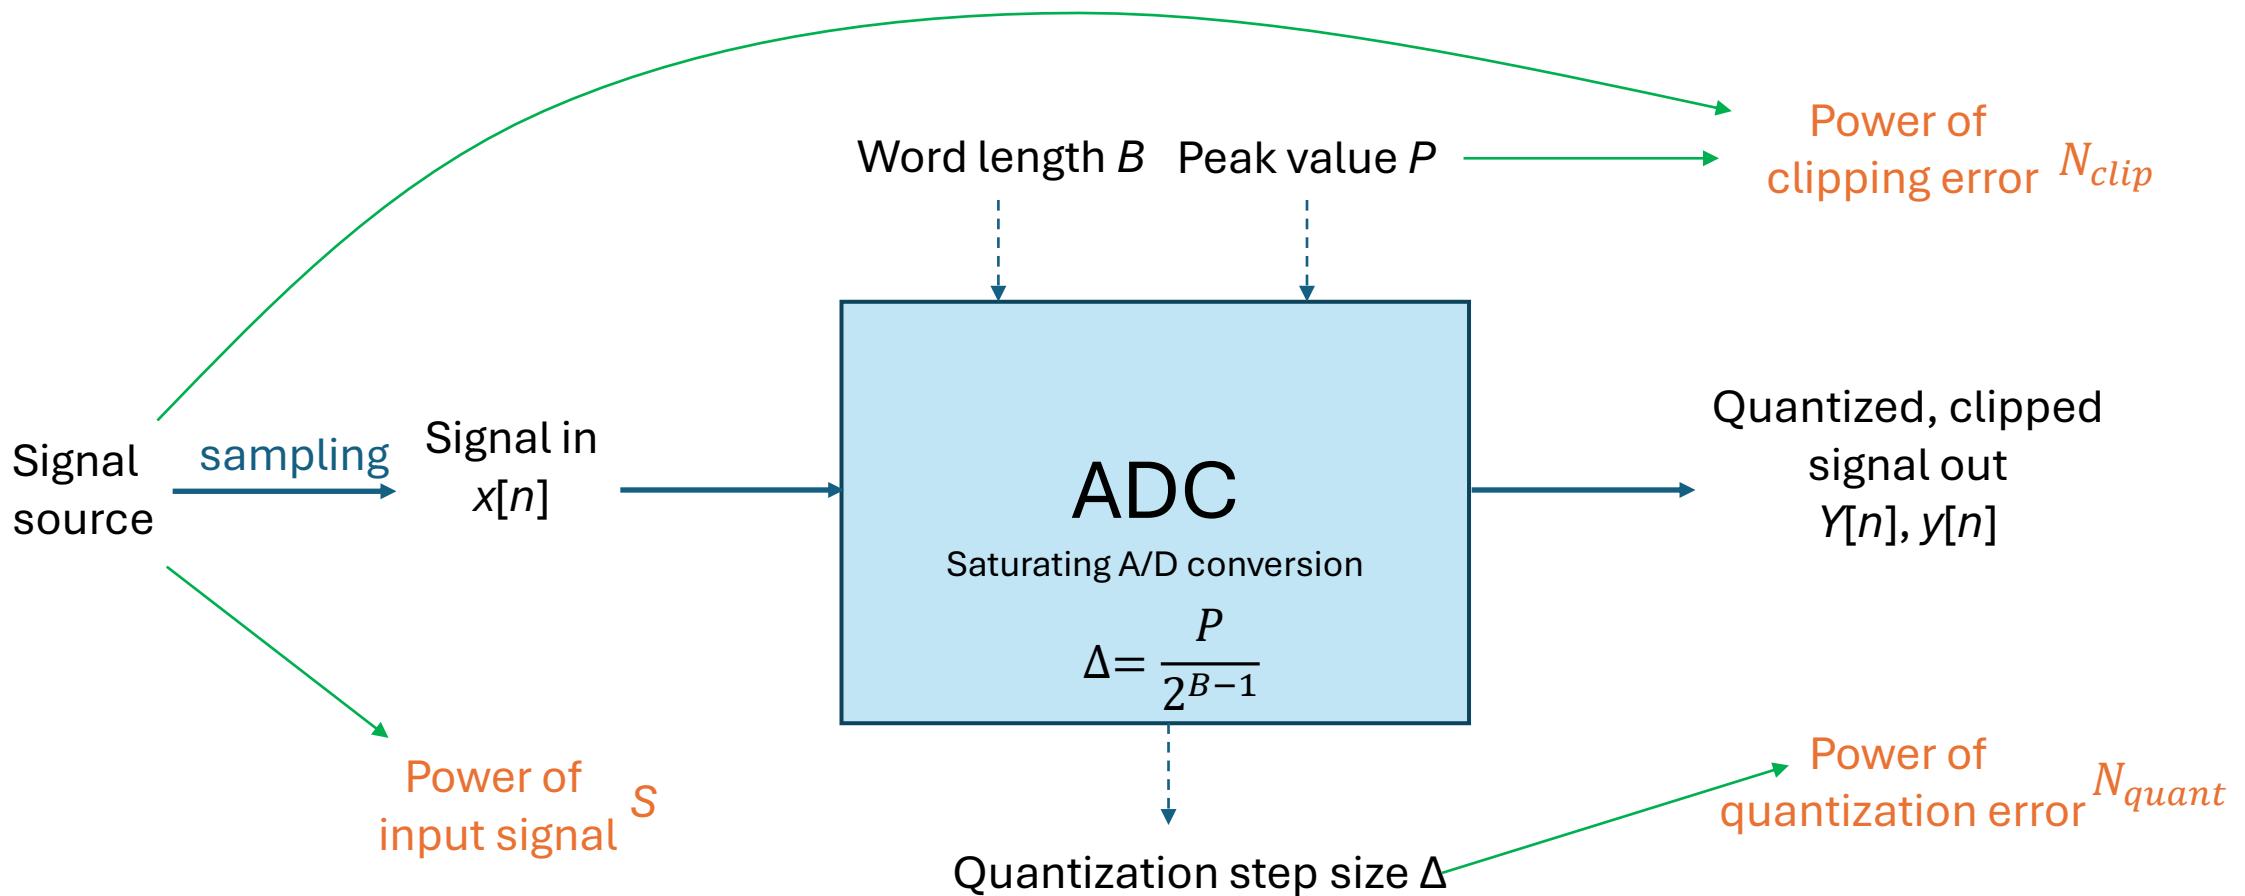
- Input signal power, S
 - Evaluated either as (1) signal source characterization or (2) from given input samples $x[n]$
- Quantization error power, N_{quant}
 - Evaluated either from (1) statistical characterization i.e. uniform distribution over $[-\Delta/2, +\Delta/2]$ or (2) from given unclipped output samples $y[n]$ and corresponding input $x[n]$
- Clipping error power, N_{clip}
 - Computed using (1) signal source characterization and saturating peak P or (2) from given clipped output samples $y[n]$ and corresponding input $x[n]$

Two ways of evaluating signal powers in each case:

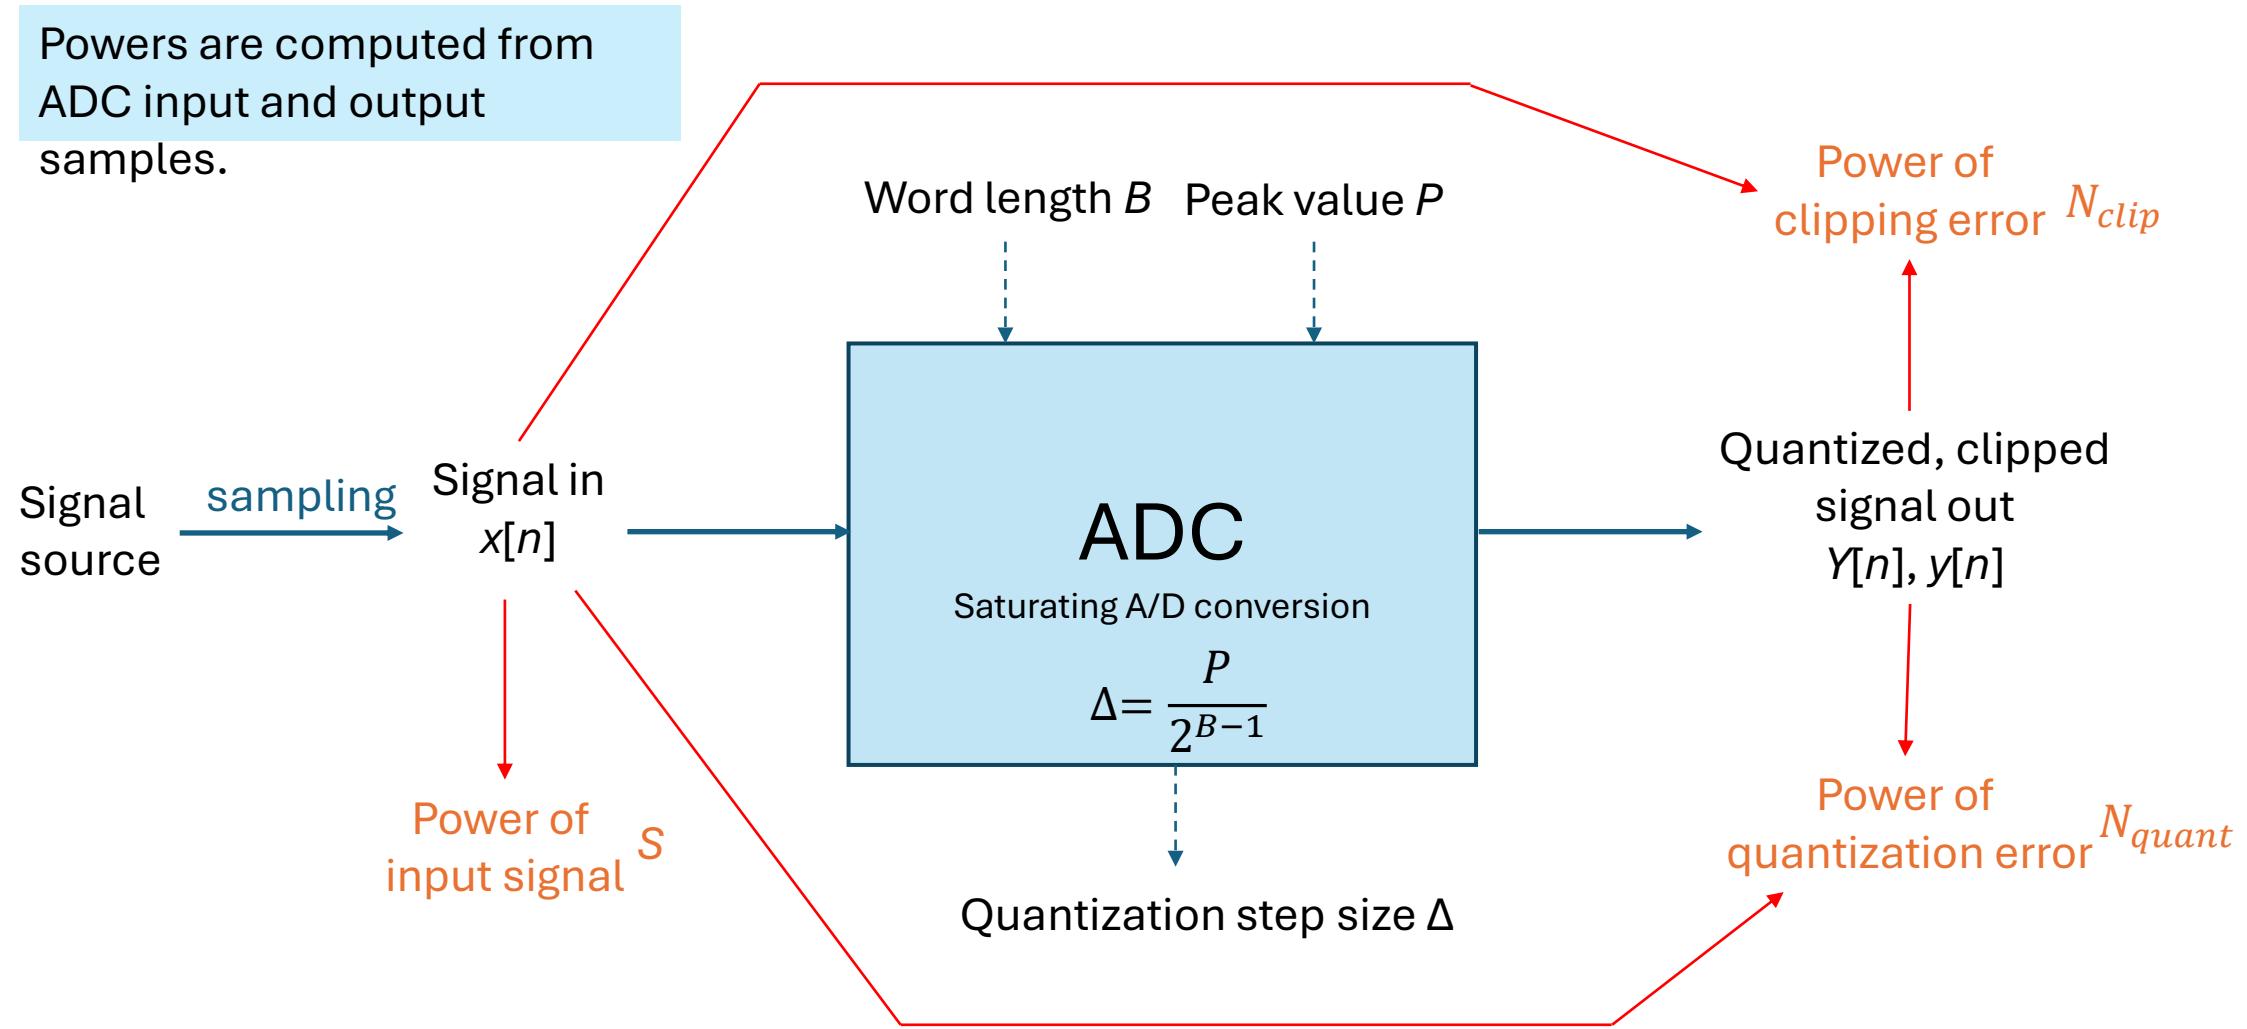
- (1) model-based *analysis*
- (2) sample-based

(1) Powers: model-based computation

Powers are computed from mathematical / statistical characterizations of the process.



(2) Powers: sample-based computation



Example. Gaussian input

(1) From model:

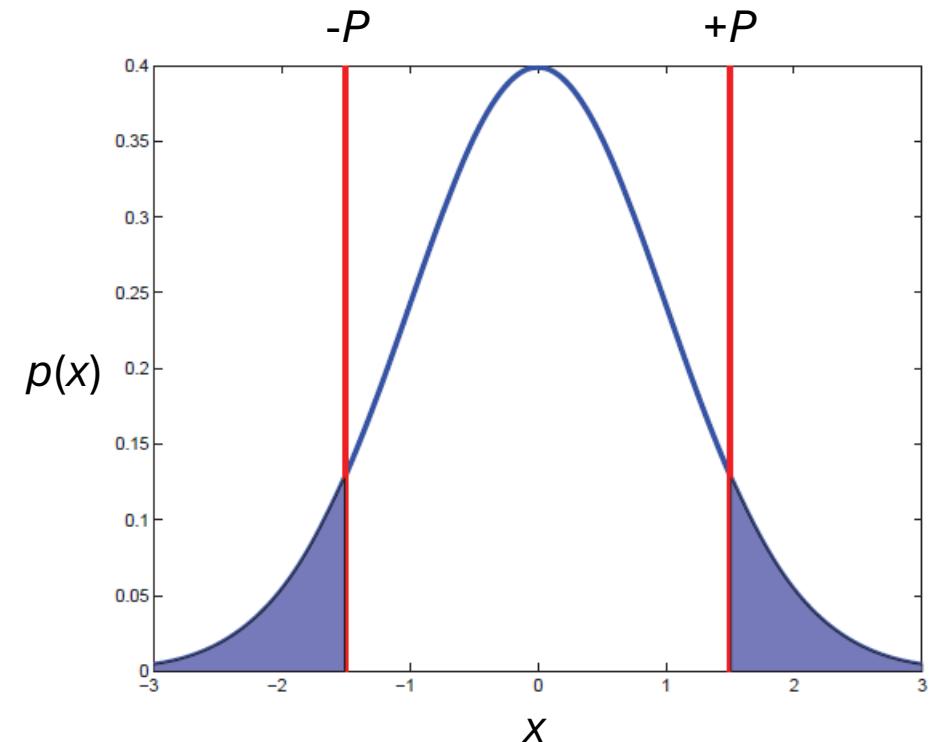
Signal source characterized by its probability density function (pdf)

Signal power $S = \sigma^2$ Variance

Clipping error power $N_{clip} = S - \int_{-P}^{+P} x^2 p(x) dx$

Quantization error power $N_{quant} = \frac{1}{\Delta} \int_{-\Delta/2}^{+\Delta/2} x^2 dx = \frac{\Delta^2}{12}$

$$\Delta = \frac{P}{2^{B-1}}$$



(2) From samples:

$$S = \frac{1}{N} \sum x[n]^2$$

$$N_{clip} = \frac{1}{N} \sum_{n \in clip} (y[n] - x[n])^2$$

$$N_{quant} = \frac{1}{N} \sum_{n \in noclip} (y[n] - x[n])^2$$

1.2. Crest factor (CF)

- Property of the input signal, defined as

$$CF_{SIG} = 10 \log_{10} \frac{Q^2}{S}$$

Q : peak value of the signal
(or value exceeded with some low probability)
 S : input signal power = RMS² (RMS = root mean square)

- Quantifies how large deviations from RMS value there can be at the ADC input
- ADC configuration must tolerate those deviations
 - Adjust saturating limit P properly
 - Then, no clips or just few

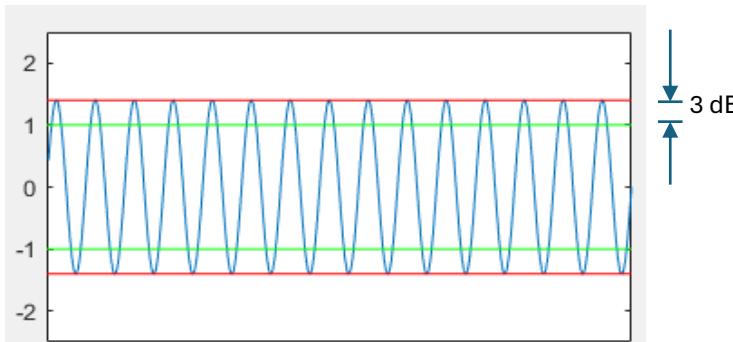
Crest factor examples

$$CF_{SIG} = 10 \log_{10} \frac{Q^2}{S} [\text{dB}]$$

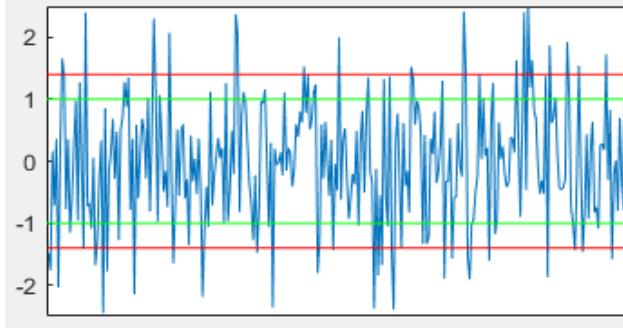
Q : peak value of the signal
(or value exceeded with some low probability)

S : input signal power = RMS² (RMS = root mean square)

Sinusoid



Gaussian signal



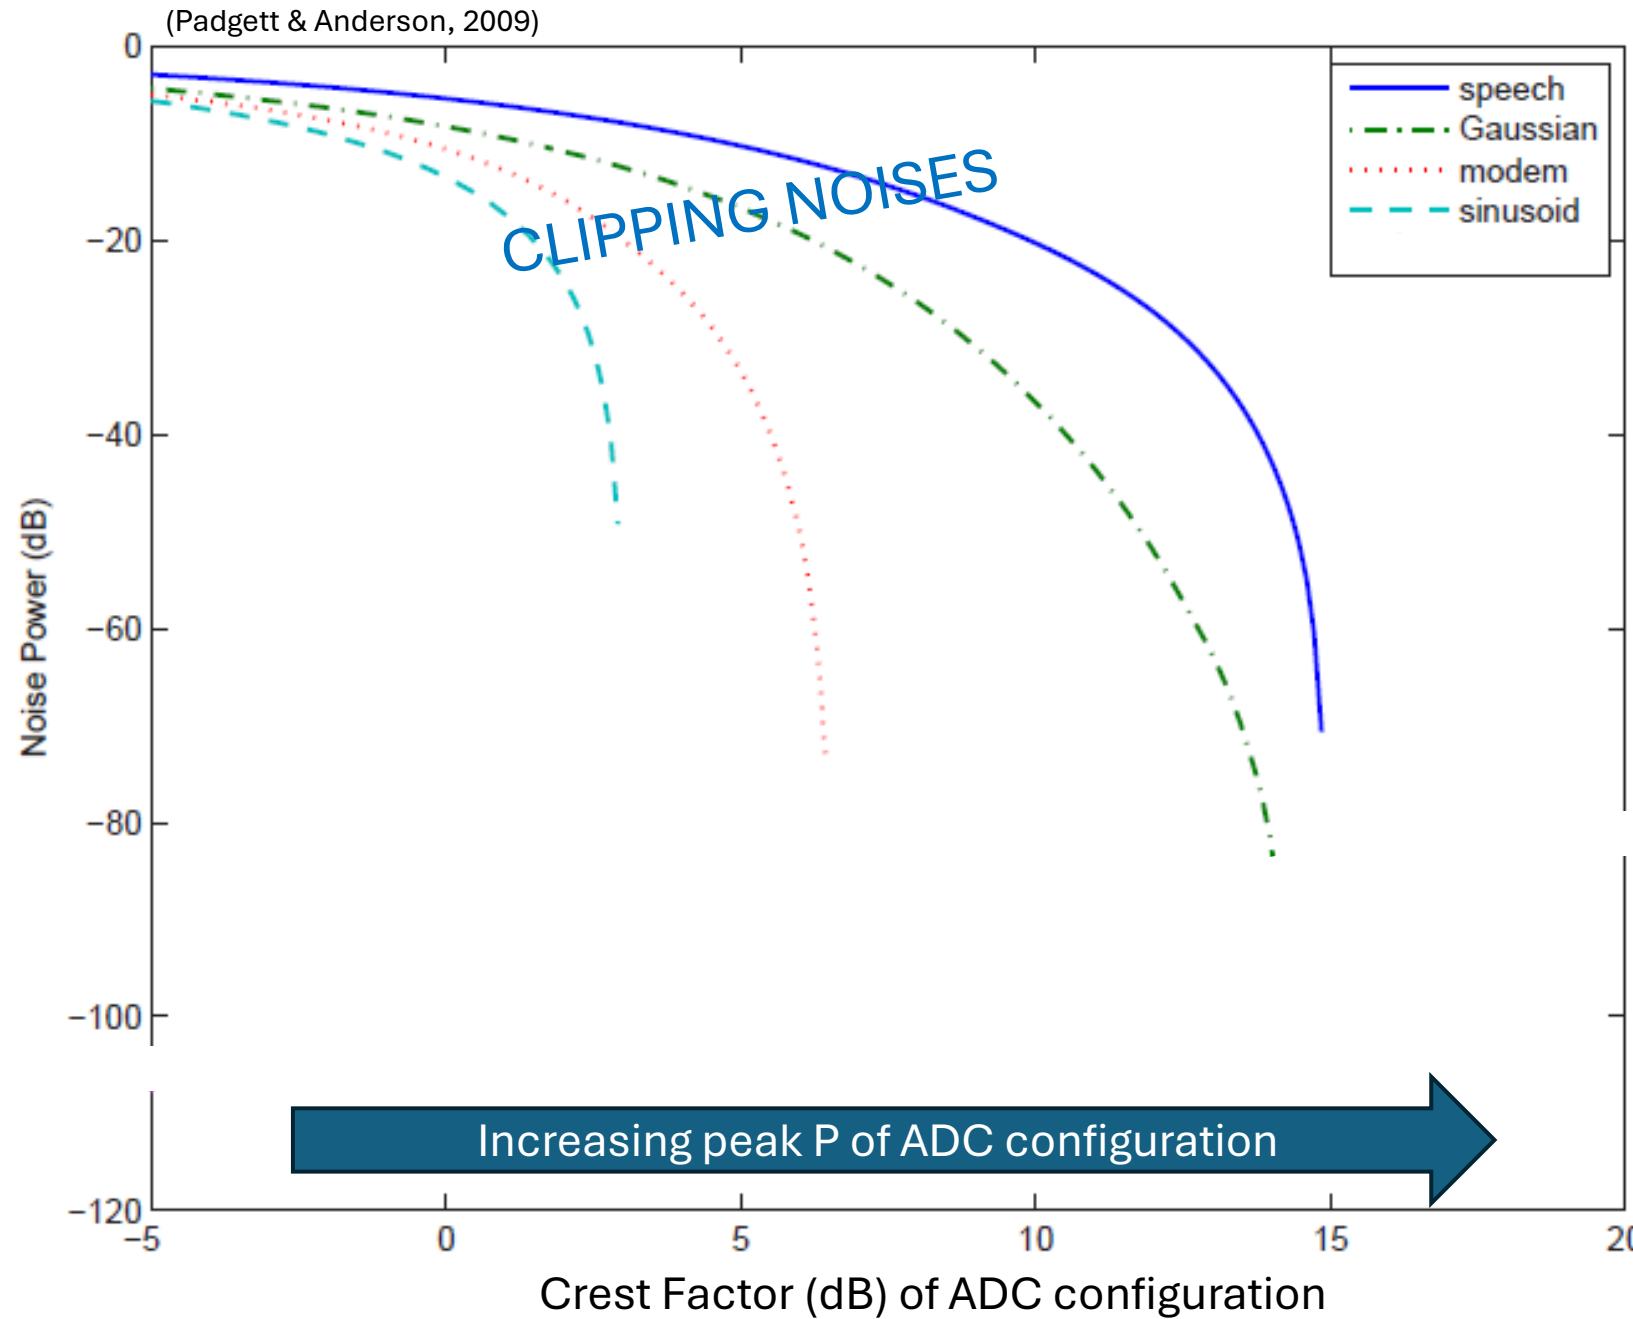
$$\text{RMS} = \frac{Q}{\sqrt{2}}$$

$$CF_{SIG} = 10 \log \frac{Q^2}{Q^2/2} = 10 \log 2 = 3.01 \text{ dB}$$

Event $|x[n]| > 3.29\sigma$ occurs with probability 0.001.

Set Q = 3.29σ in the equation =>

$$CF_{SIG} = 10 \log \frac{(3.29\sigma)^2}{\sigma^2} = 20 \log 3.29 = 10.34 \text{ dB}$$

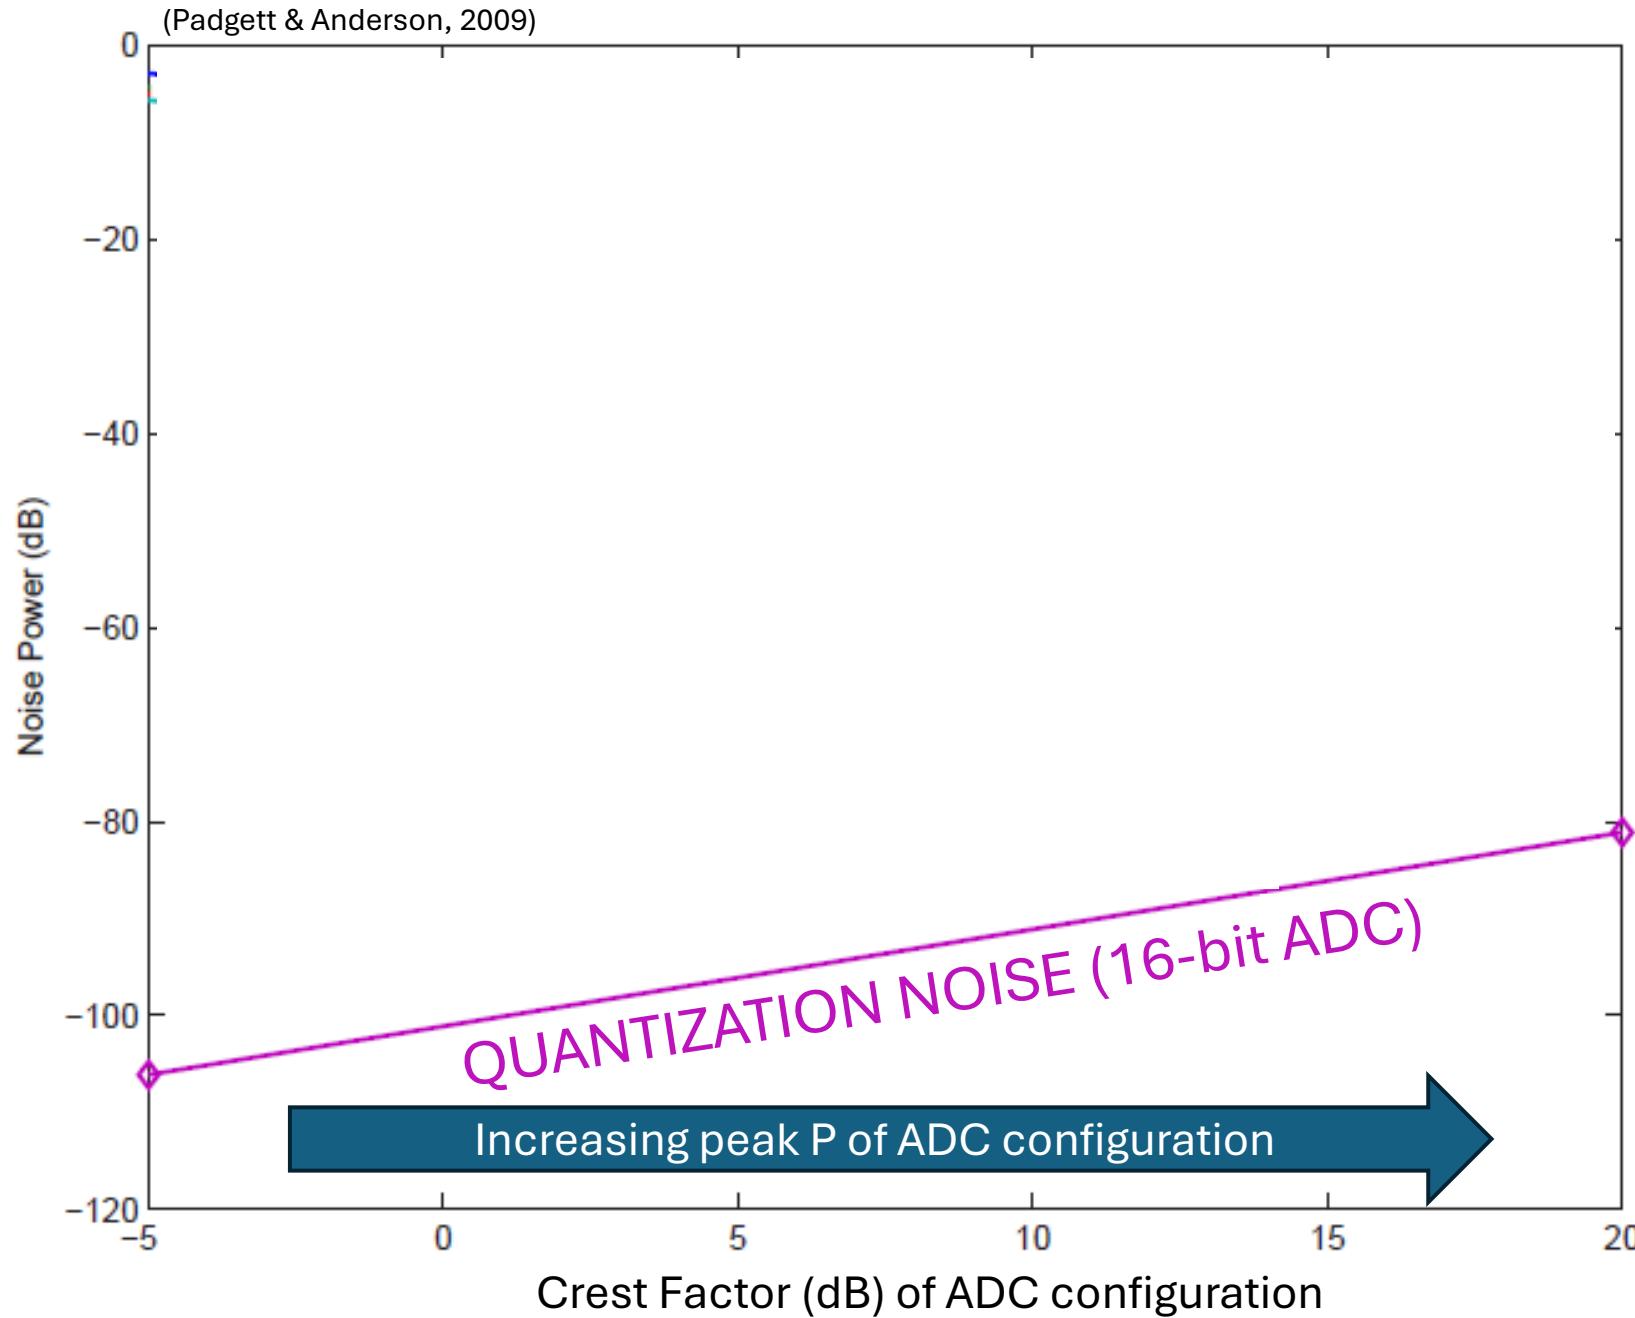


With increasing peak P level clipping noises drop.

It depends on the signal how much P must be increased.

Speech signal has heavier tails than Gaussian.

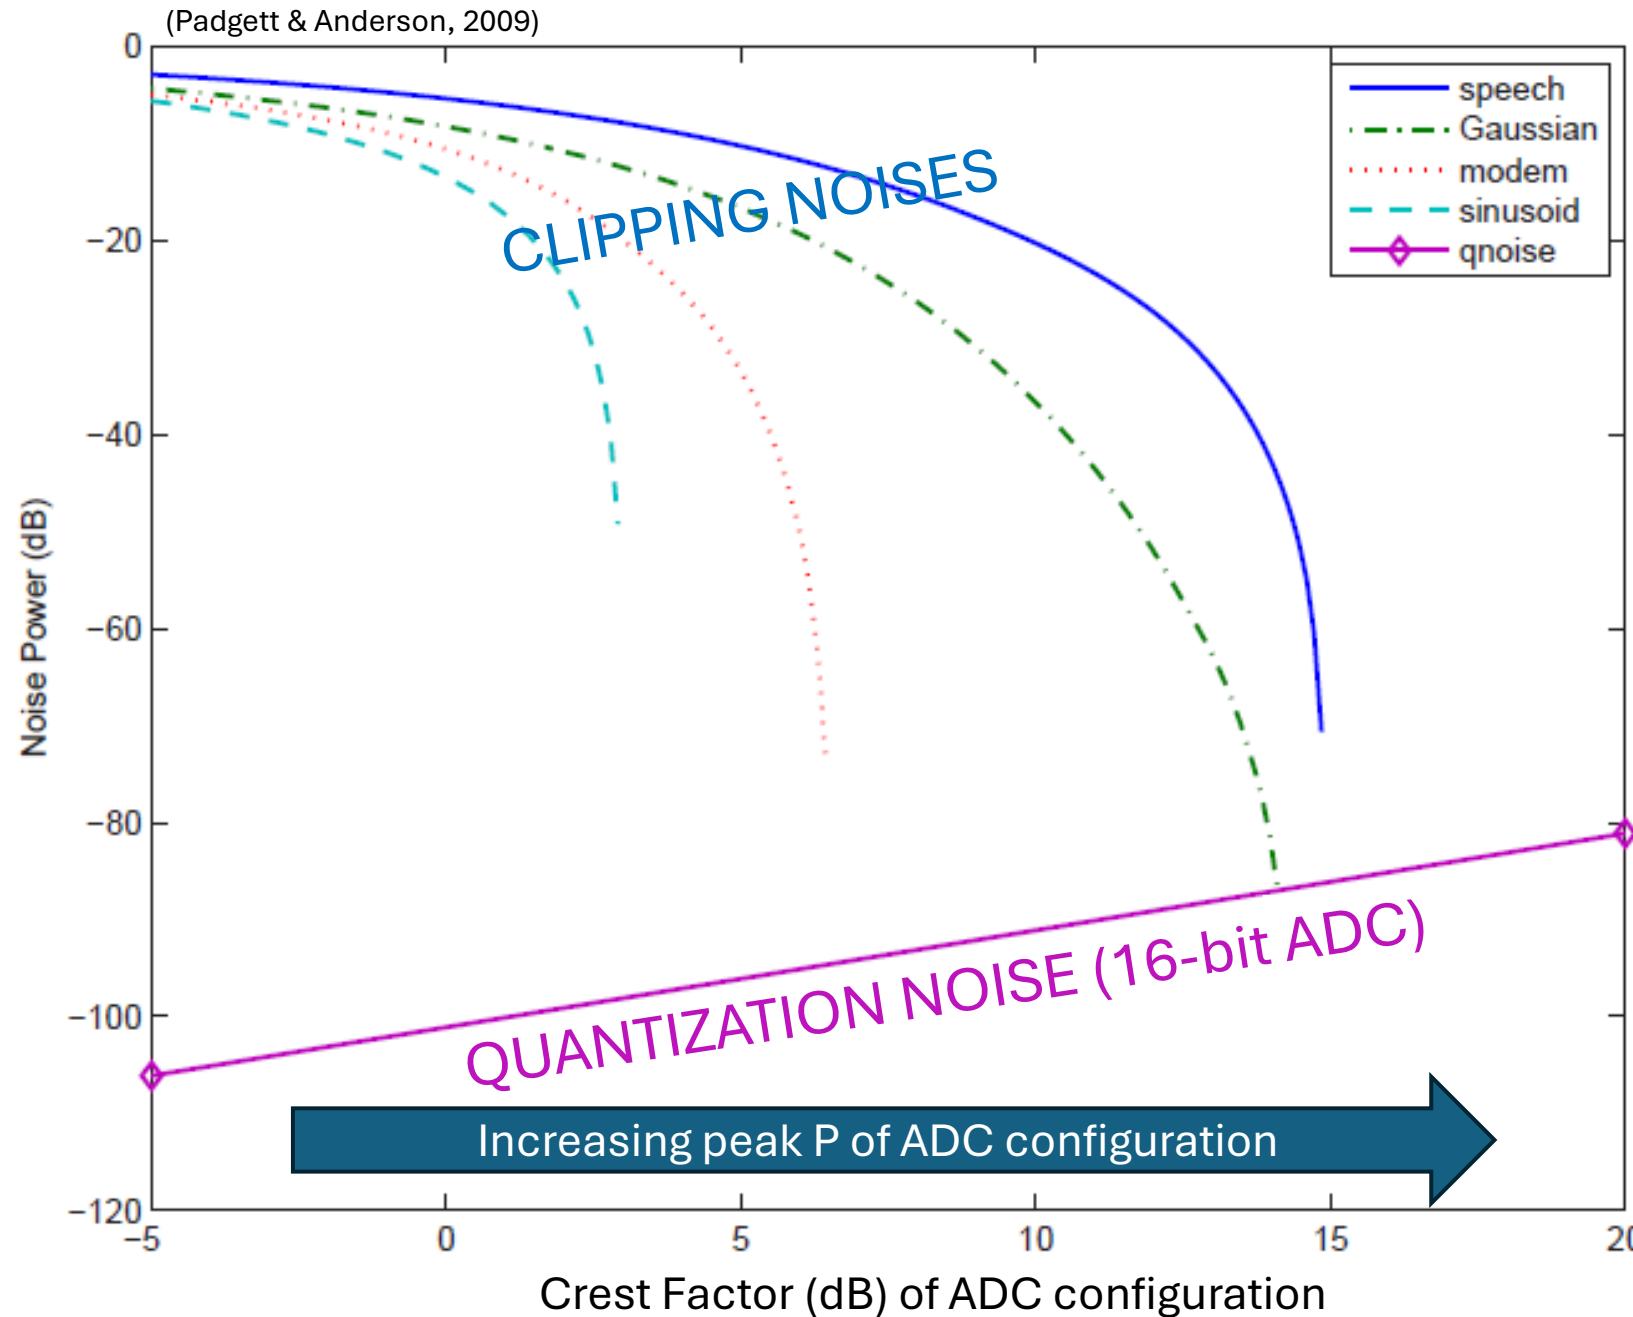
$$CF_{ADC} = 10 \log_{10} \frac{P^2}{S} [\text{dB}]$$



With increasing peak P level quantization noise increases as quantization level difference increases.

$$\Delta = \frac{P}{2^{B-1}}$$

$$CF_{ADC} = 10 \log_{10} \frac{P^2}{S} [\text{dB}]$$



Considering ADC noise,
optimum tuning of P is at
the crossing point of these
two noises.

$$CF_{ADC} = 10 \log_{10} \frac{P^2}{S} [\text{dB}]$$

Matlab example

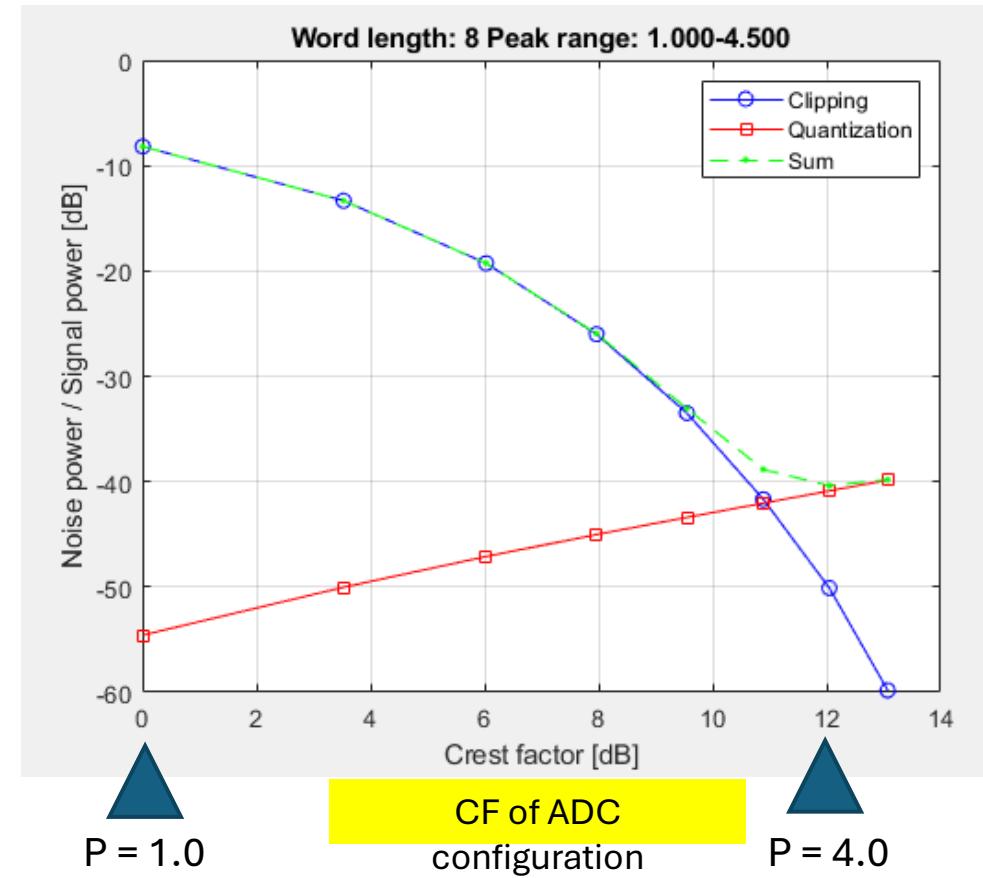
- Two Matlab files (given in Moodle):

- saturatingADC.m
- crestFactorPlot.m

```
% Gaussian signal with variance 1
x = randn(1,1000000);
Ps = 1.0:0.5:4.5; % ADC peaks P
B = 8; % ADC output word length
crestFactorPlot(x,Ps,B);
```

We see that clipping and quantization noise powers are similar when ADC is tuned for CF = 11 dB (ADC peak $P = 3.5\sigma$).

Perhaps this is a good choice for ADC peak value as quantization and clipping noises are in balance.



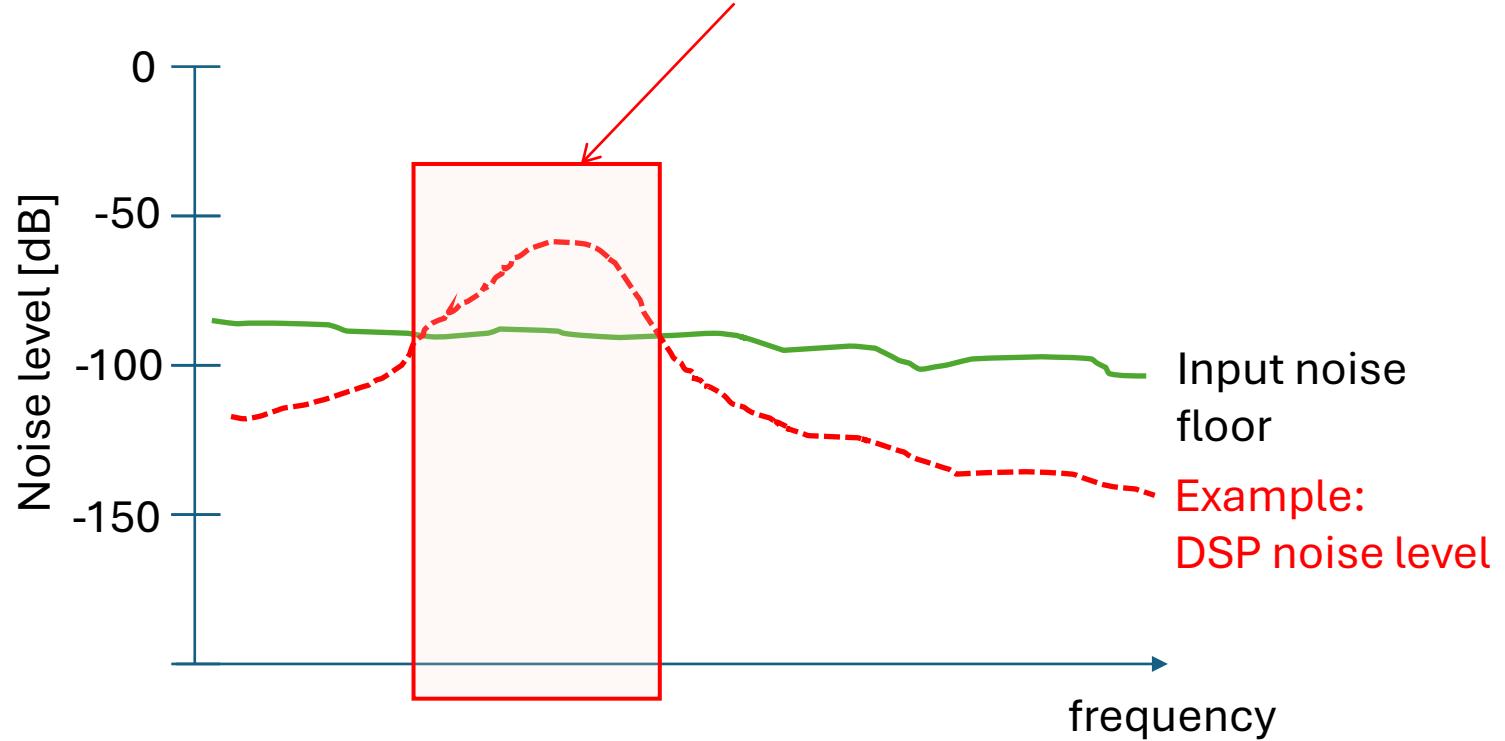
1.3. Input noise floor

The noise induced by A/D conversion and inherent noise in the analog input set up a noise floor for processing.

For the signal processing output, we cannot get above this level.

So, this input noise floor can set up a goal for DSP design: **keep the noise at the level of input noise!**

In this frequency range, the DSP noise level degrades the signal => goal for improvement!



(Ifeachor & Jervis, 2002, p. 817)

1.4. Spectrum of the quantization noise

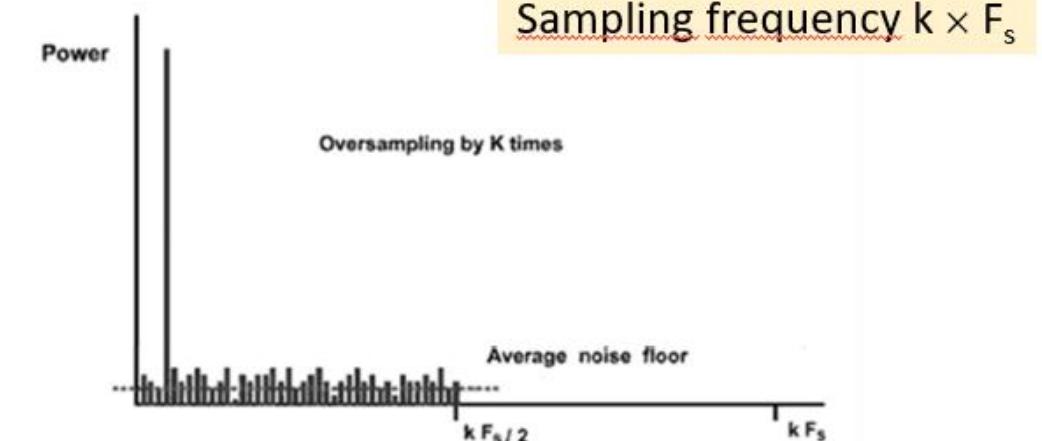
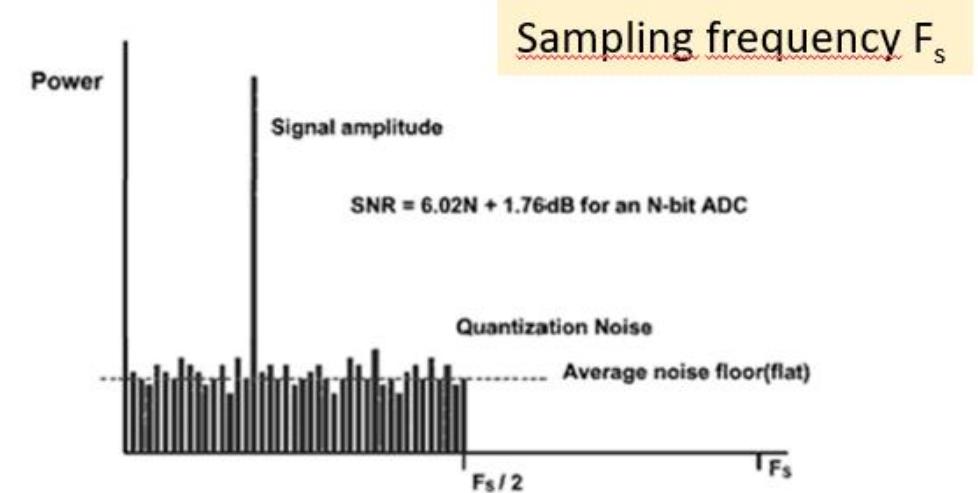
- Spectrum of the ADC quantization noise is approximately flat
- The power spectral density (PSD) for sample rate F_s is

$$S(f) = \frac{N_{quant}}{F_s}$$

where

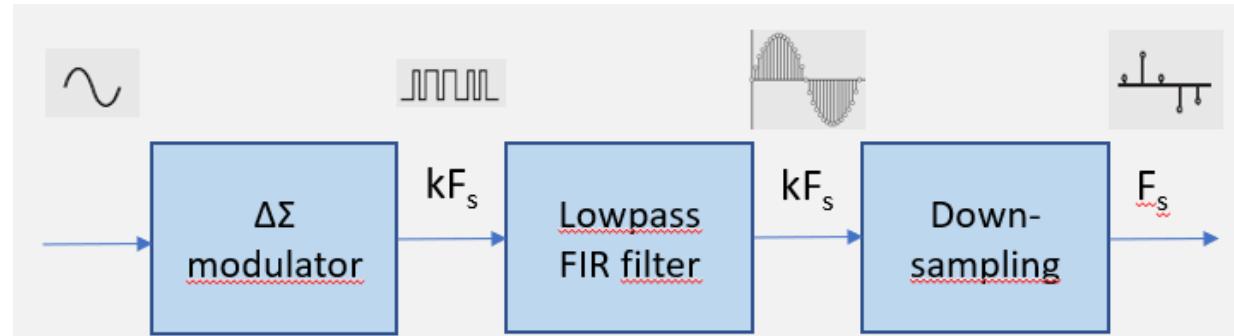
$$N_{quant} = \frac{1}{\Delta} \int_{-\Delta/2}^{+\Delta/2} x^2 dx = \frac{\Delta^2}{12}$$

- **Note:** If sampling rate is increased, noise floor becomes lower

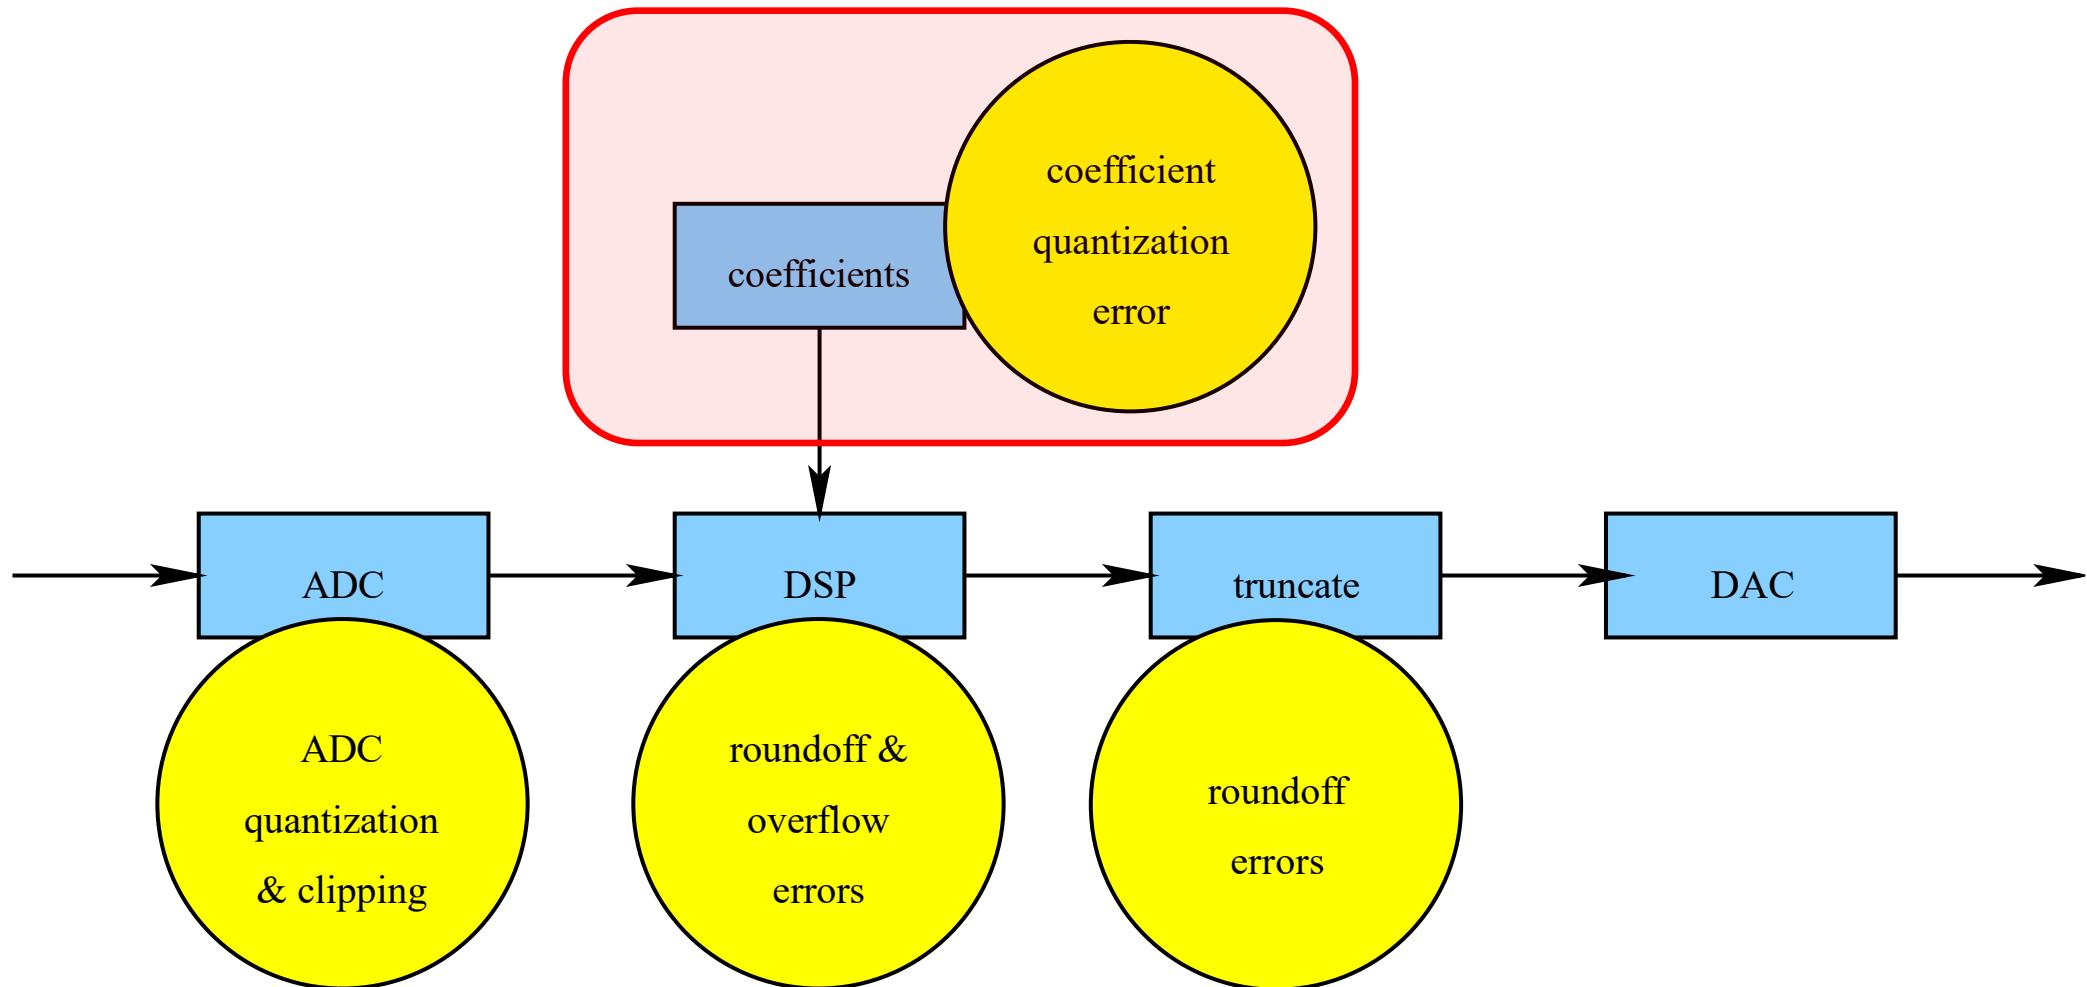


Oversampling ADC

- Exploit this notion about sample rate and noise floor
- Oversampling ADC is a combination of
 - High-frequency sampling
 - A/D conversion
 - Decimation to lower the sample rate
- Extreme: One can even use one-bit ADC with very high sampling rate
 - $\Delta\Sigma$ modulator
- Result: ADC quantization noise reduced
- Common technique in audio signal processing
- A multirate technique: more on this later



2. IIR filter coefficient quantization

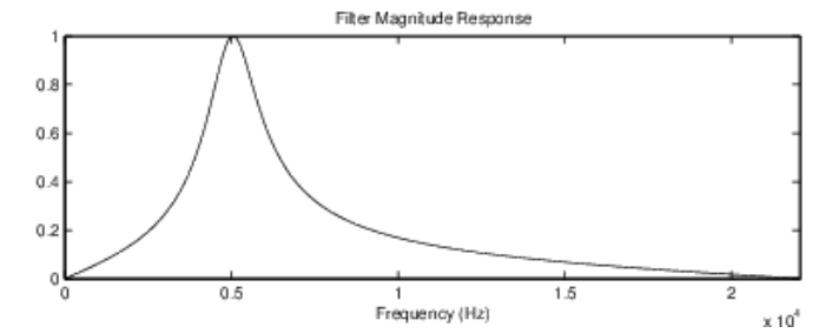
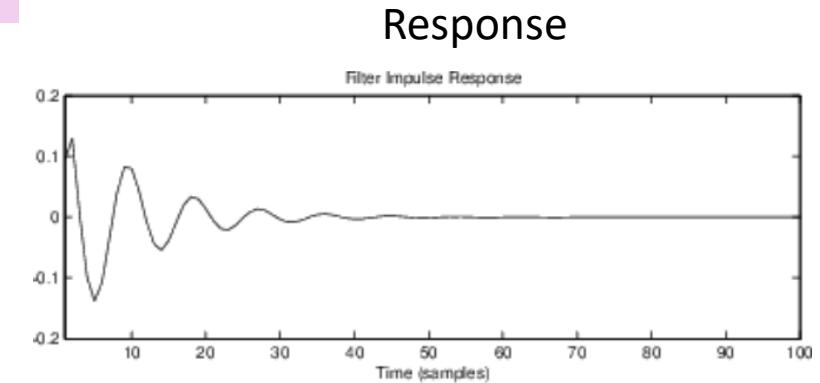
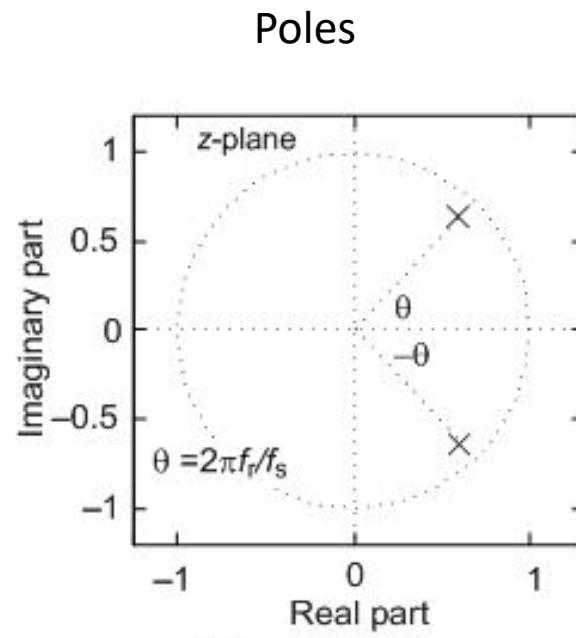
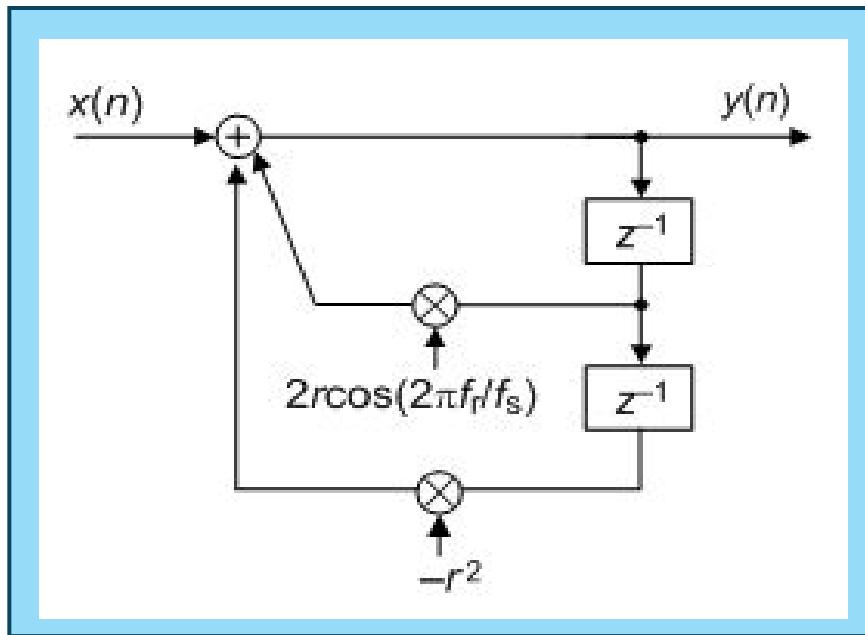


Problem

- When we quantize IIR filter coefficients
 - What happens to the response of the filter?
 - Is the filter still stable?
- How to find optimal solutions?
 - Not just coefficient quantization, also choice of implementation structure important
- Considering two cases
 - 2nd order IIR resonator
 - High-order IIR filter

Case 1: 2nd order IIR resonator

- Resonator: impulse response of the filter is a damped sinusoid with frequency f_r and damping factor r
- Is a band-pass filter

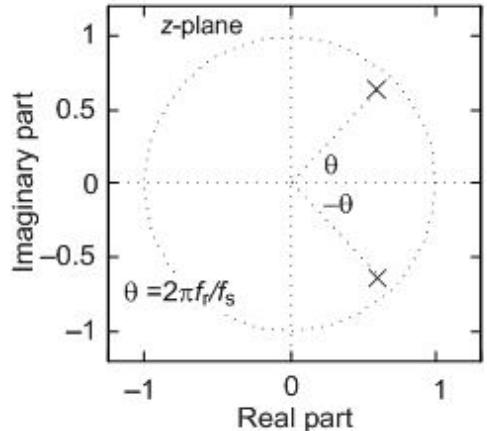


2nd order IIR resonator

- The transfer function of the filter is of form

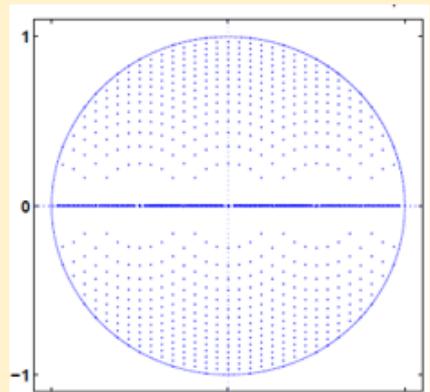
$$H(z) = \frac{1}{1-a_1z^{-1}-a_2z^{-2}}$$

- The filter has complex-conjugate roots (poles) at $z = re^{\pm j\theta}$, where $\theta = 2\pi f_r/f_s$
- In direct fixed-point implementation, we need to quantize $a_1 = 2r \cos \theta$ and $a_2 = -r^2$
 - But: not full control of pole positions
 - Observable sparsities in pole position plots
 - Especially low-pass (f_r close to 0) and high-pass (f_r close to $f_s/2$) filters are harder to realize - longer word lengths needed

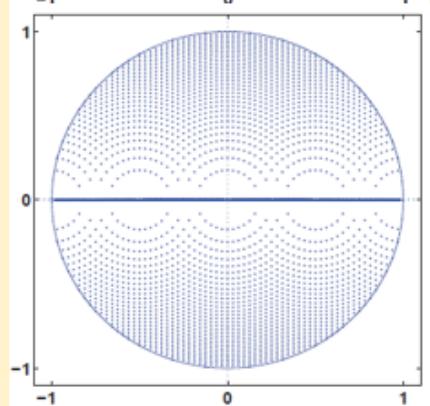


Available pole positions

5-bit coefficients



6-bit coefficients

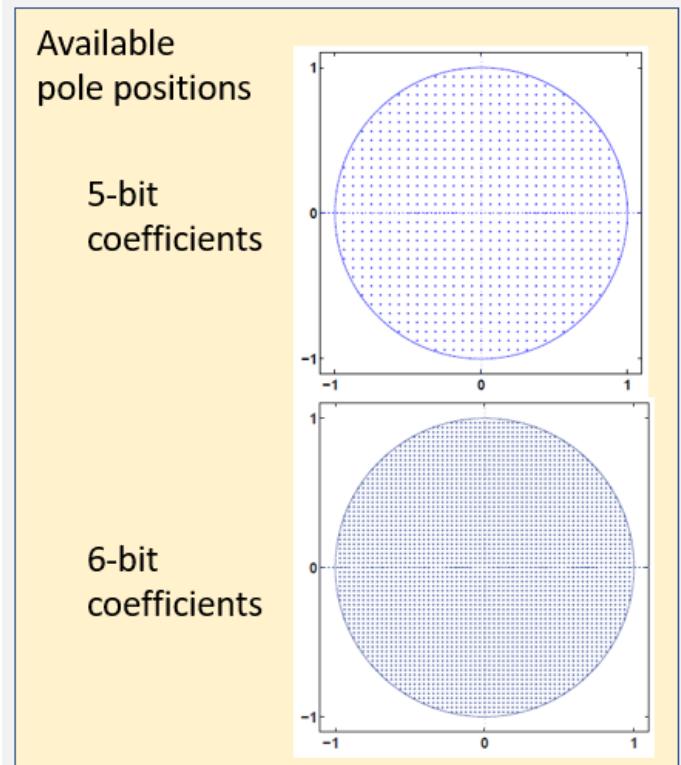
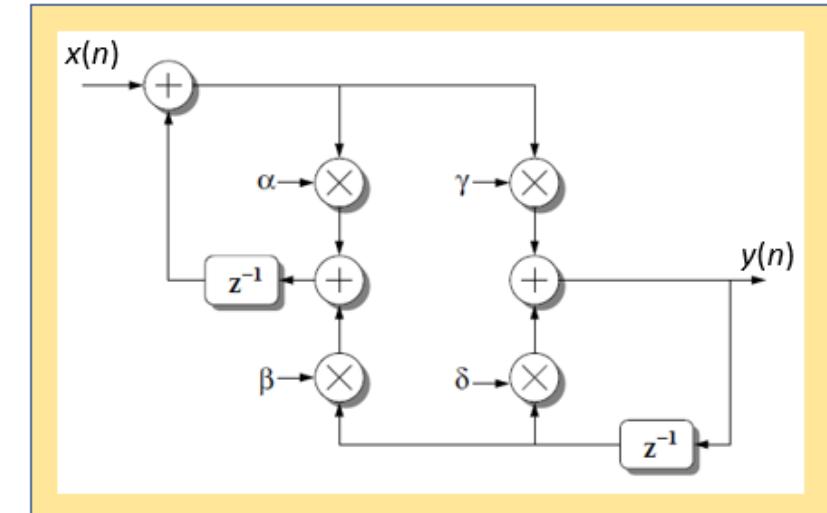


2nd order coupled form IIR

- Provides a solution to resonator implementation problems
 - note: four multiplications instead of two
- The transfer function is

$$H(z) = \frac{\gamma}{1-(\alpha+\delta)z^{-1}-(\beta\gamma-\alpha\delta)z^{-2}}$$

- Setting $\alpha = \delta = r \cos \theta$ and $\beta = -\gamma = r \sin \theta$, we get resonator's transfer function (up to scaling γ)
- Quantization of coefficients ($\alpha, \beta, \gamma, \delta$) leads to **rectangular grid** of pole positions

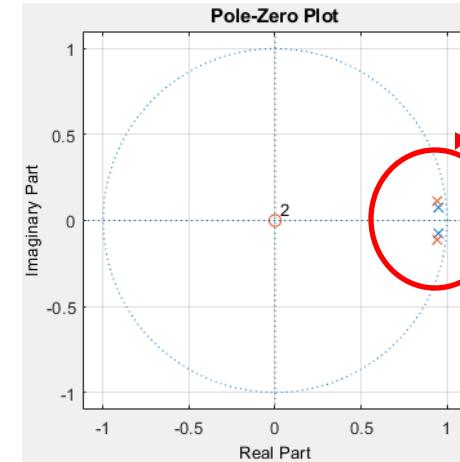
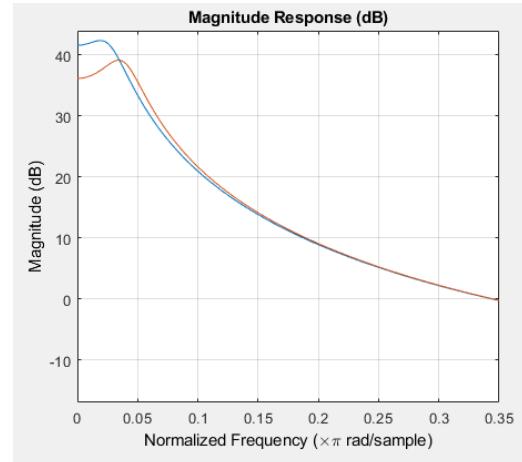


Discussed in <https://www.dsprelated.com/showarticle/183.php>

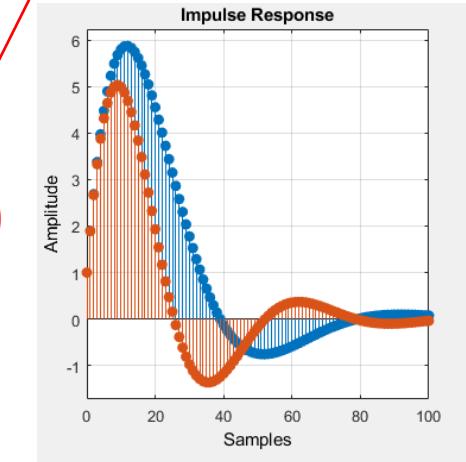
Example. Resonator for $f_r = f_s/80$, $r = 0.95$

$\theta = 0.025\pi$.

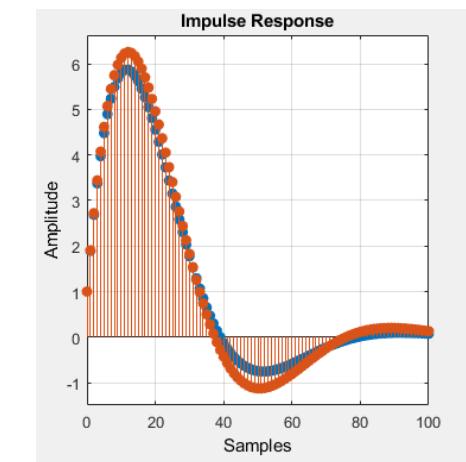
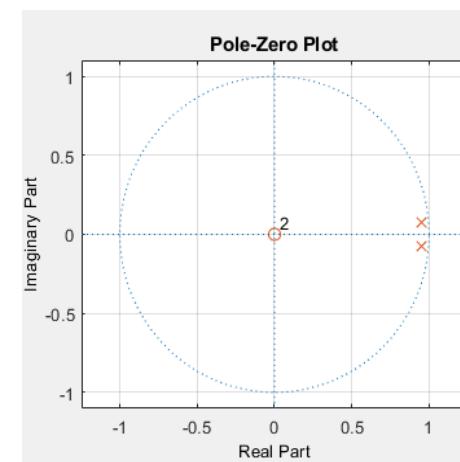
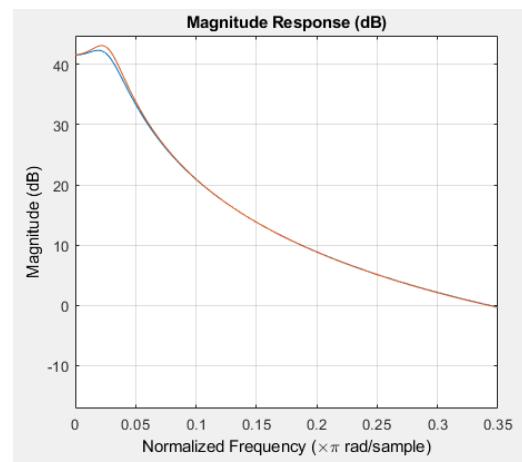
Direct form fixed-point
s8.6 format



difference



Coupled form fixed-point
s8.6 format

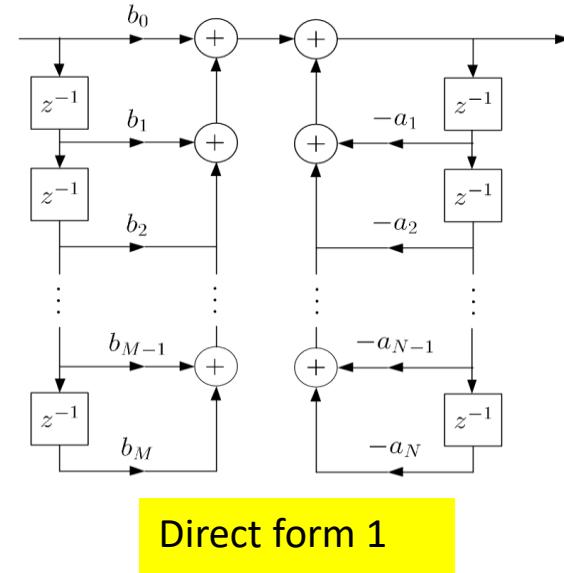


— floating-point filter

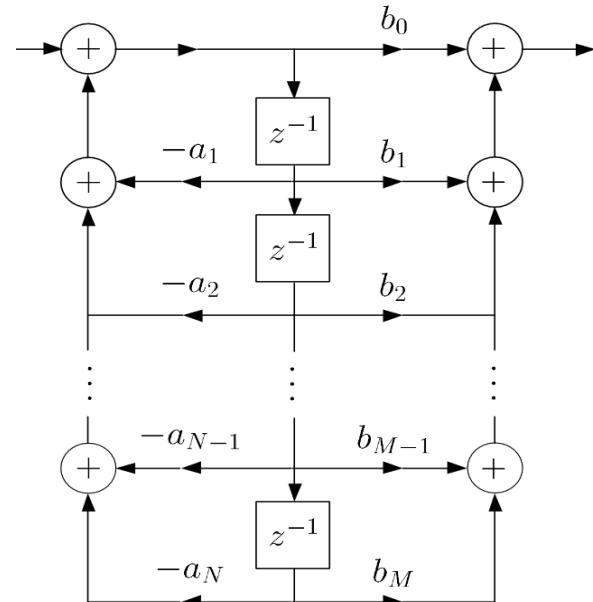
— fixed-point filter

Case 2: high-order IIR filters

- Direct form implementations are problematic. Reason:
 - Typically, we want to use the same fixed-point format for all numerator coefficients and another one for all denominator coefficients
 - However, the **magnitude of the coefficients may vary a lot**, and then the precision of some coefficients can become too low



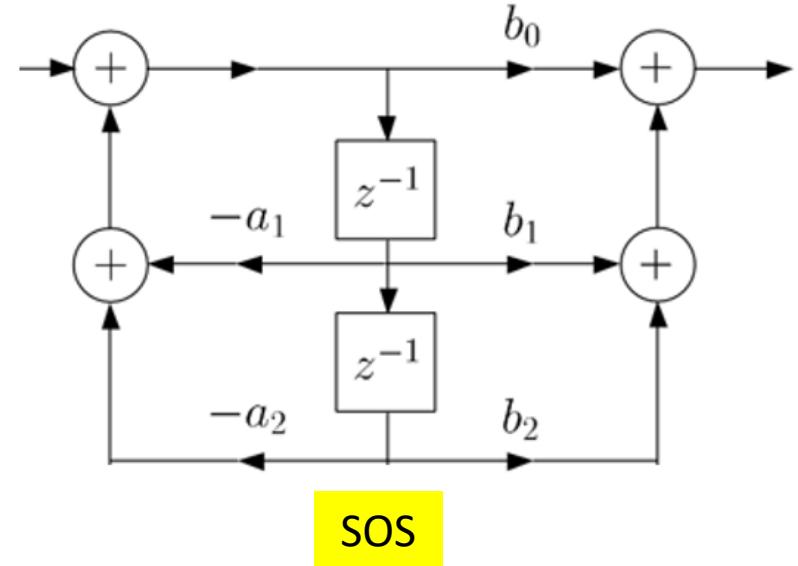
Direct form 1



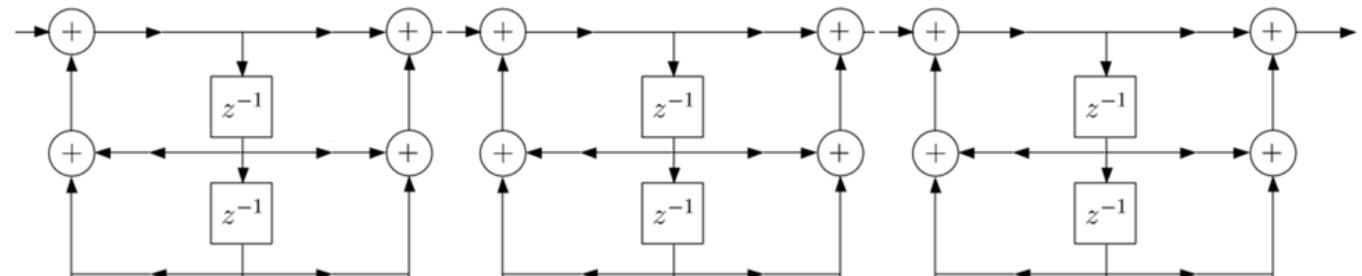
Direct form 2 (canonic)

Case 2: high-order IIR filters

- The solution is to consider cascaded second-order section (SOS) implementations
 - Coefficients of each section have relatively small range
 - Note: not first-order sections as we want to avoid complex arithmetic – 2nd order SOS has real coefficients



SOS



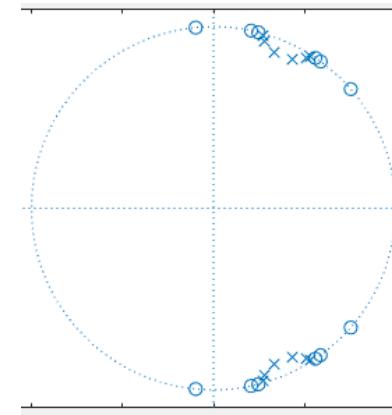
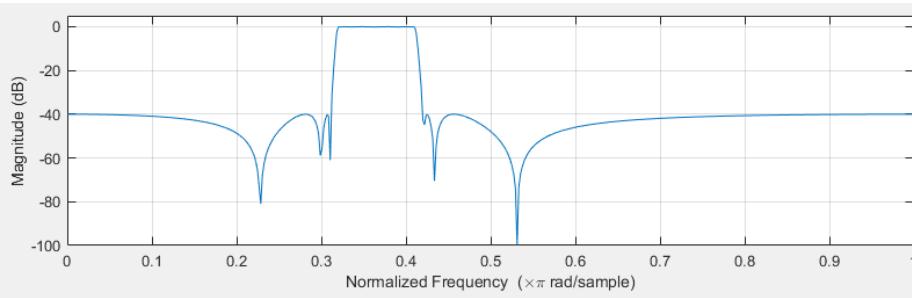
SOS cascade

Example. elliptic bandpass filter (order 12)

Specification

- passband cut-offs 0.32, 0.41 ($\times F_s/2$)
- passband ripple 0.1 dB
- stopband attenuation 40 dB

Property: fast transition in gain



Designed using matlab: `[b, a] = ellip(6, 0.1, 40, [0.32, 0.41])`



numerator
 b [0...12]

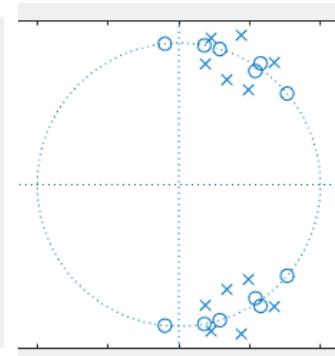
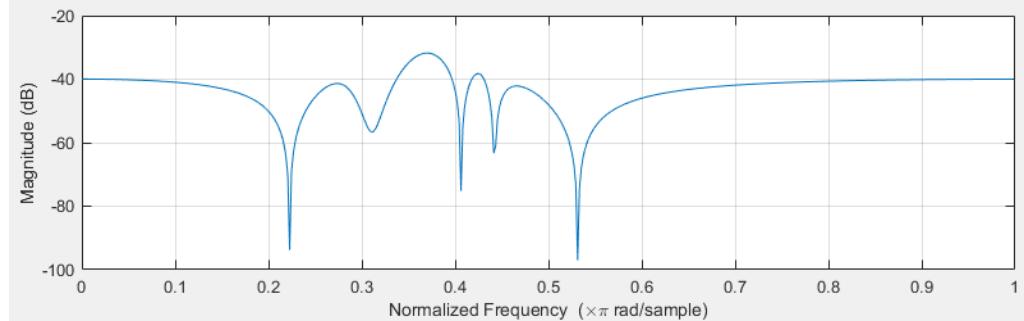
```
0.010402294826284
-0.046981738177133
0.140584418862936
-0.291587543984709
0.483154936264528
-0.638421915199626
0.705123004845763
-0.638421915199625
0.483154936264528
-0.291587543984708
0.140584418862936
-0.046981738177132
0.010402294826284
```

denominator
 a [0...12]

```
1.000000000000000
-4.736891331870121
14.770074008149525
-31.339008666596676
52.290097782458098
-68.277884768136573
73.209425975402013
-63.111029848304227
44.673344222037109
-24.741797451330484
10.774944914870538
-3.191627411828423
0.622743578239246
```

Large variation in magnitude

Trying 16-bit fixed-point quantization, one format for a and one for b , that is, s16.15 for b and s16.8 for a



No passband, not even stable!

Solution: Cascaded SOS implementation

Matlab:

```
[b,a] = ellip(6,0.1,40,[0.32,0.41])  
[SOS,G] = tf2sos(b,a)
```

6 SOS sections

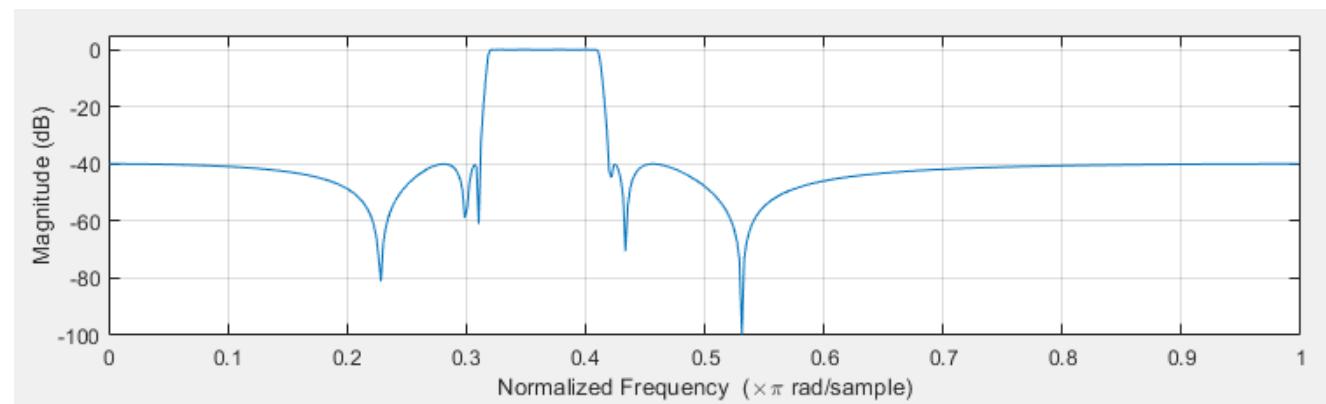
$b[0]$	$b[1]$	$b[2]$	$a[0]$	$a[1]$	$a[2]$
1	0.195430137972181	1.000000000000039	1	-0.664513985463847	0.852588727845048
1	-1.508636017182940	0.999999999999994	1	-0.873355320444148	0.860012128003494
1	-0.412640684596677	0.999999999998724	1	-0.564681165798558	0.932247834363709
1	-1.177232208175213	1.000000000000585	1	-1.019134004433838	0.939497814911048
1	-0.492675670575205	1.000000000001395	1	-0.543644167472686	0.983728242794488
1	-1.120723792458221	0.999999999999252	1	-1.071562688257045	0.985740948417509

Gain G: 0.010402294826284

denominator

Magnitudes are of the same order
(but note: separated gain G)

Trying 16-bit fixed-point quantization, one format for a and b , s16.14. Gain quantized separately as s16.21

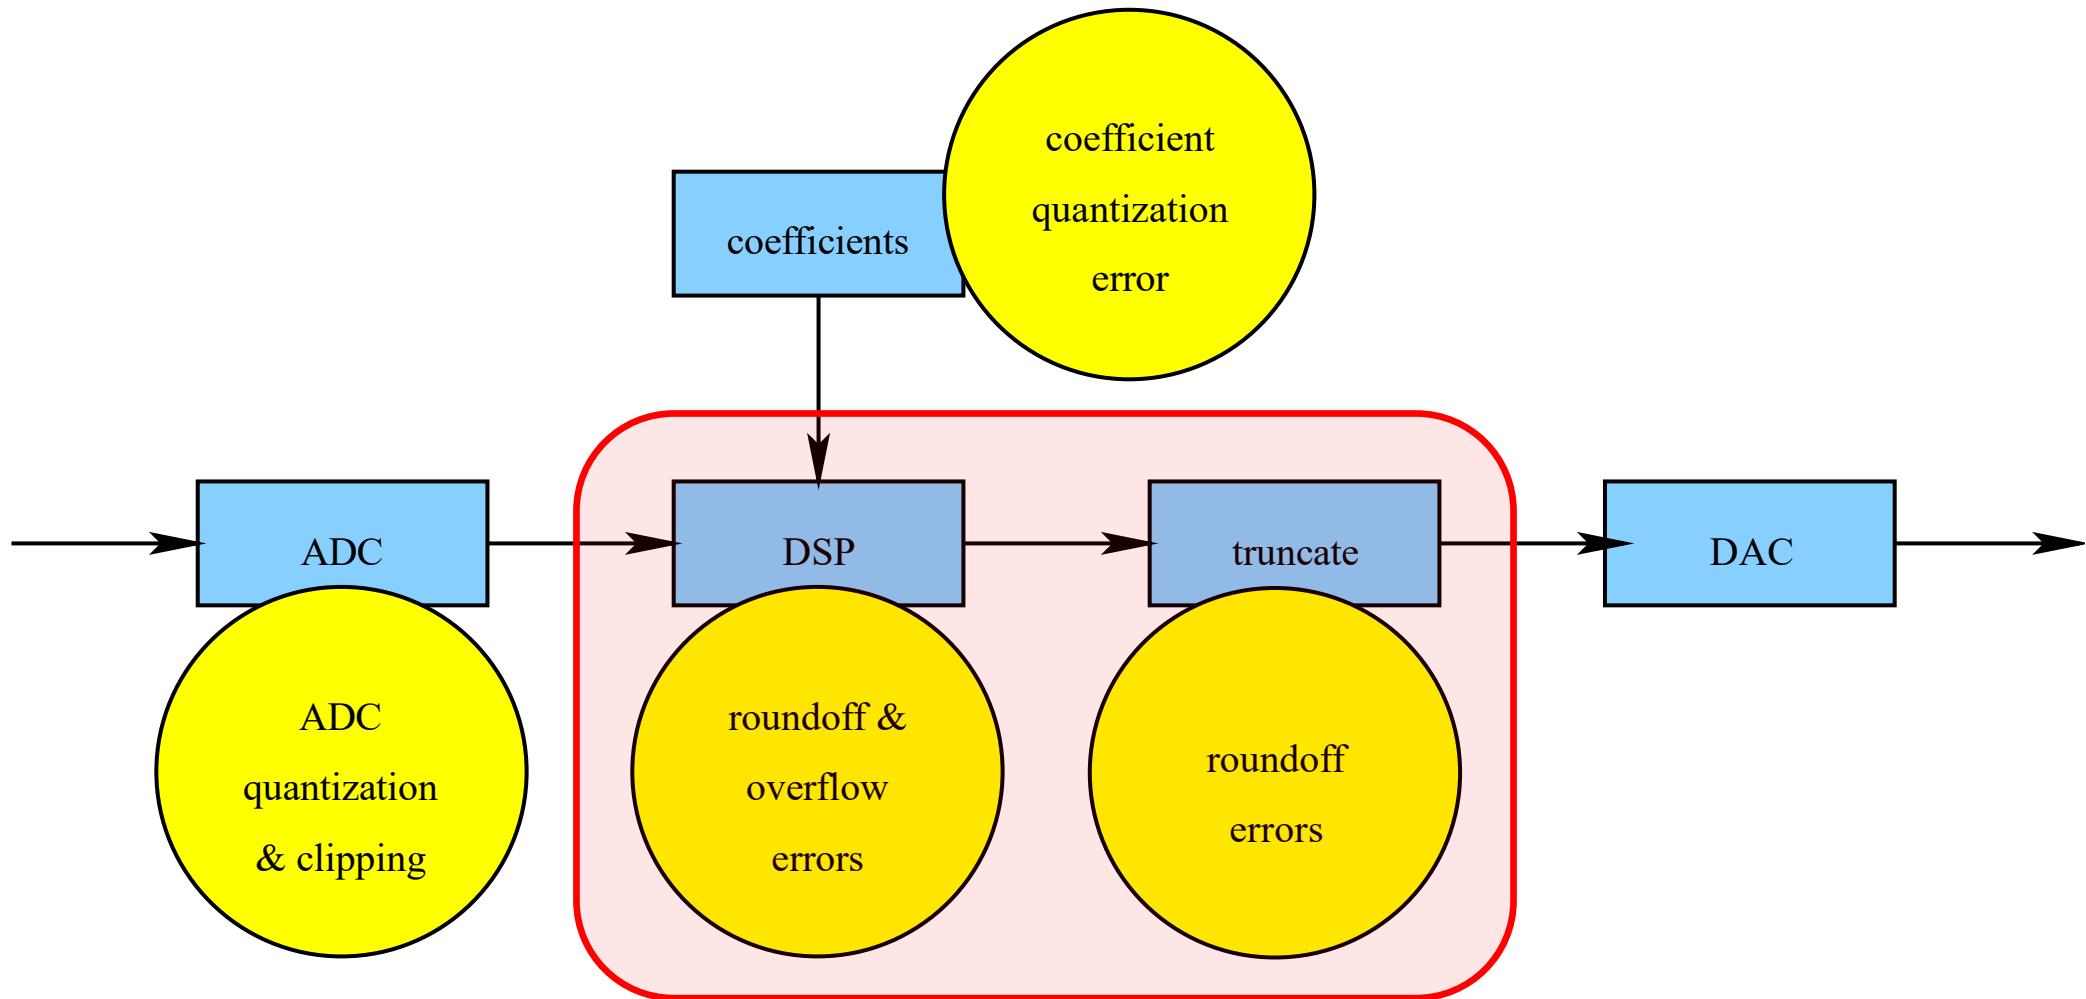


Very close to the ideal floating-point response 😊

Lesson

- To deal with coefficient quantization effects, fixed-point IIR filter design typically requires **special flow of computations**
 - It's not just coefficient quantization
 - Problem is due to limited dynamic range of fixed-point numbers
- Another example of having special structures for fixed-point computation
 - recall matrix inversion and use of decomposition techniques (see Lecture 3)

3. Quantized IIR filter: signal processing error



In the following

Implementing fixed-point IIR filter simulation takes effort.

To grasp some ideas, let us consider implementation of second-order-section (SOS) IIR filter on a DSP platform.

- Word length tunings in the platform possible.
- Learning how **input scaling** is used to deal with the overflow problems.
- Then, we consider **choice of fixed-point types** for different points on the data path.
- Based on this, **Matlab code for simulating** the system written.
- Finally, some simulation **experiments** to study effects

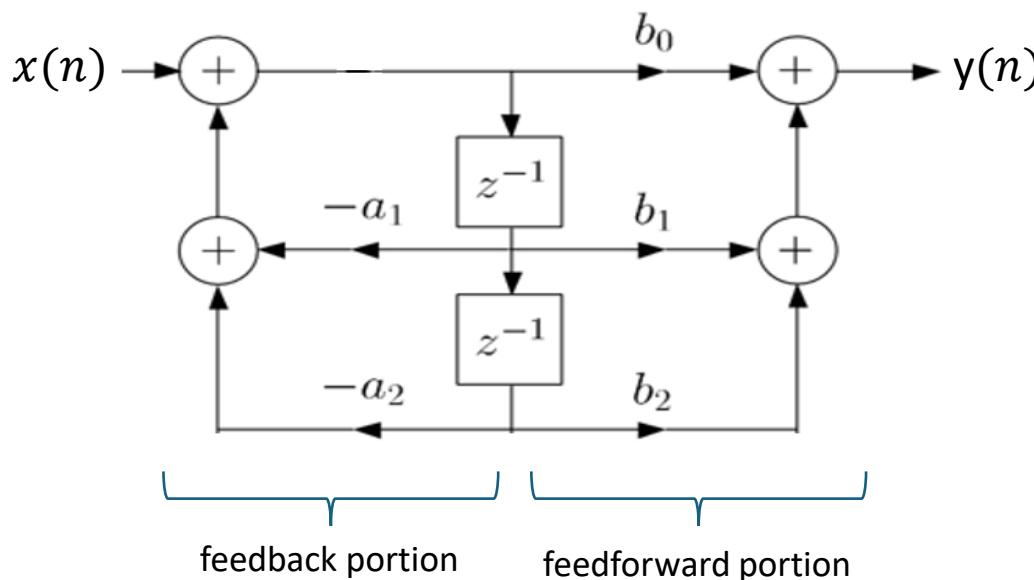
3.1. Problem: Implementing a second-order IIR filter

SOS is very useful and common component in fixed-point IIR filter implementations.

SOS transfer function:

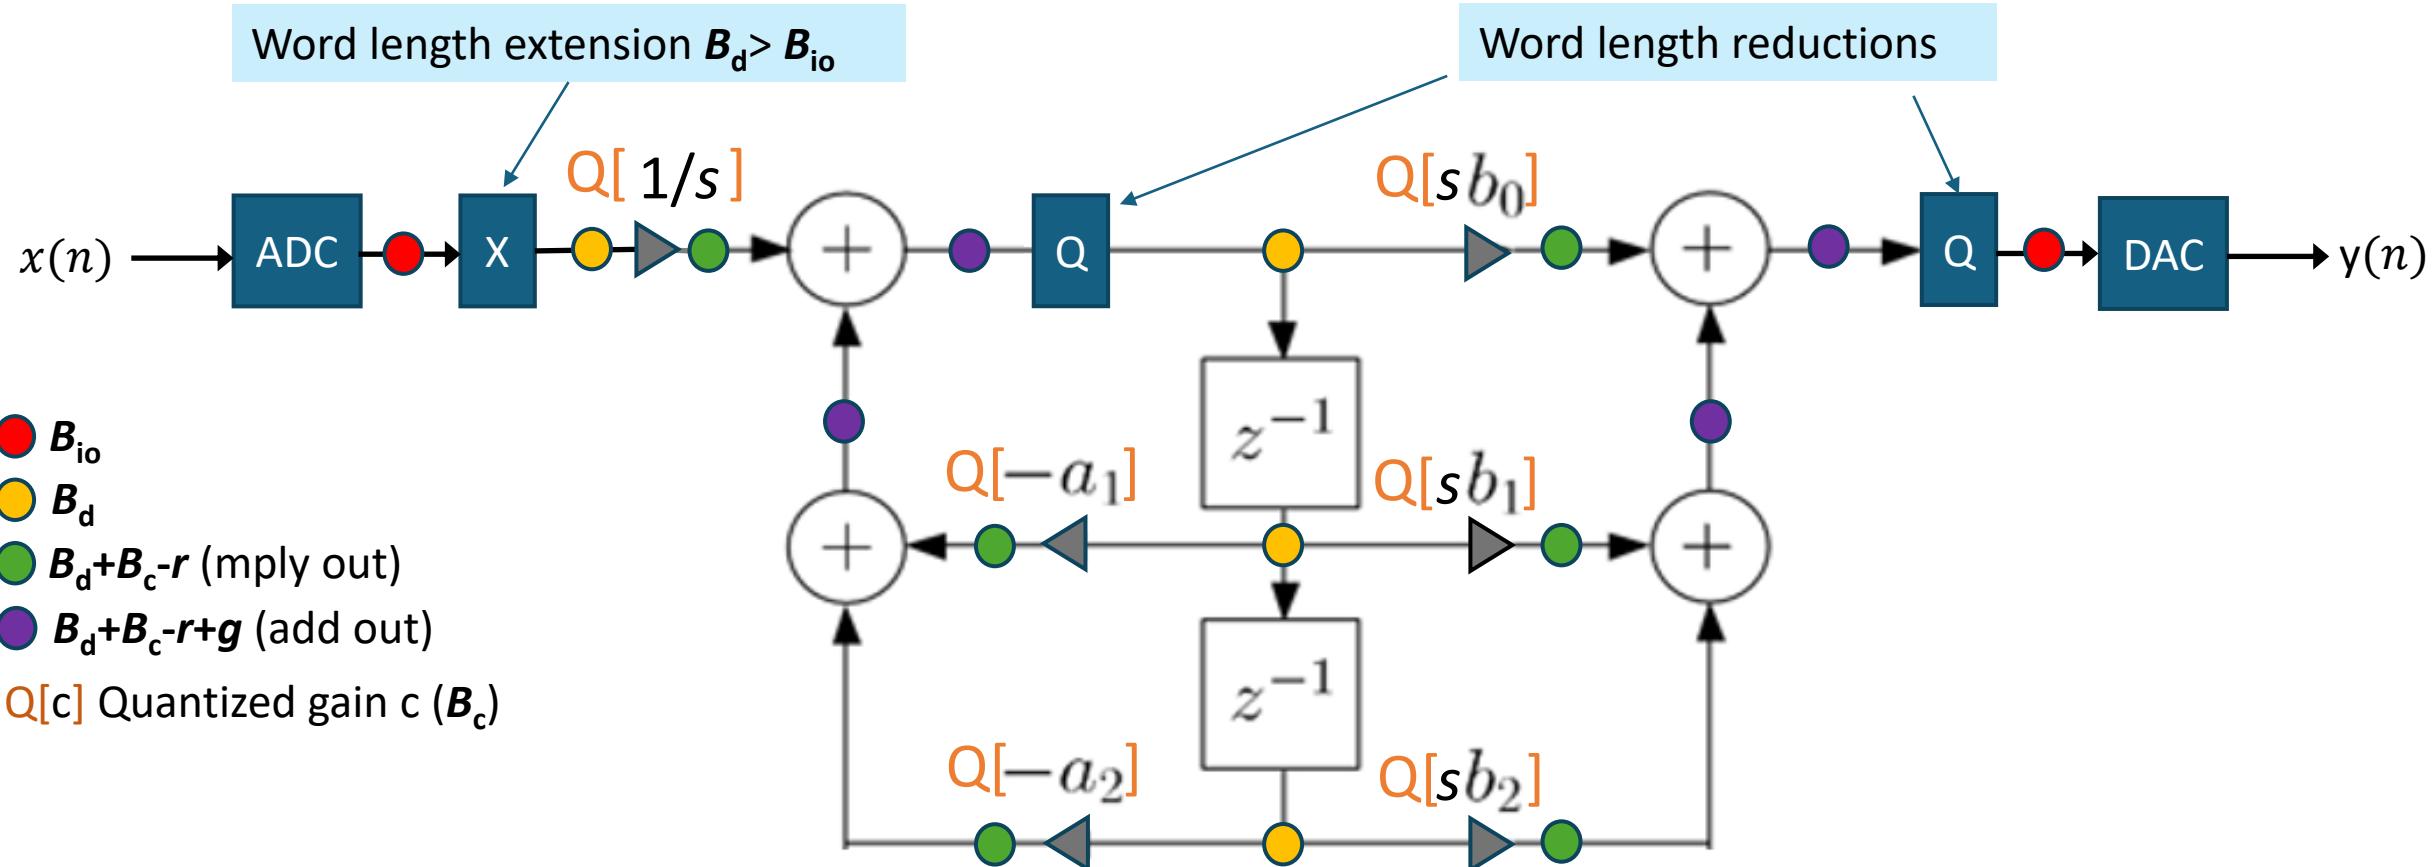
$$H(z) = \frac{b_0 + b_1 z^{-1} + b_2 z^{-2}}{1 + a_1 z^{-1} + a_2 z^{-2}}$$

Canonic direct form II –based implementation

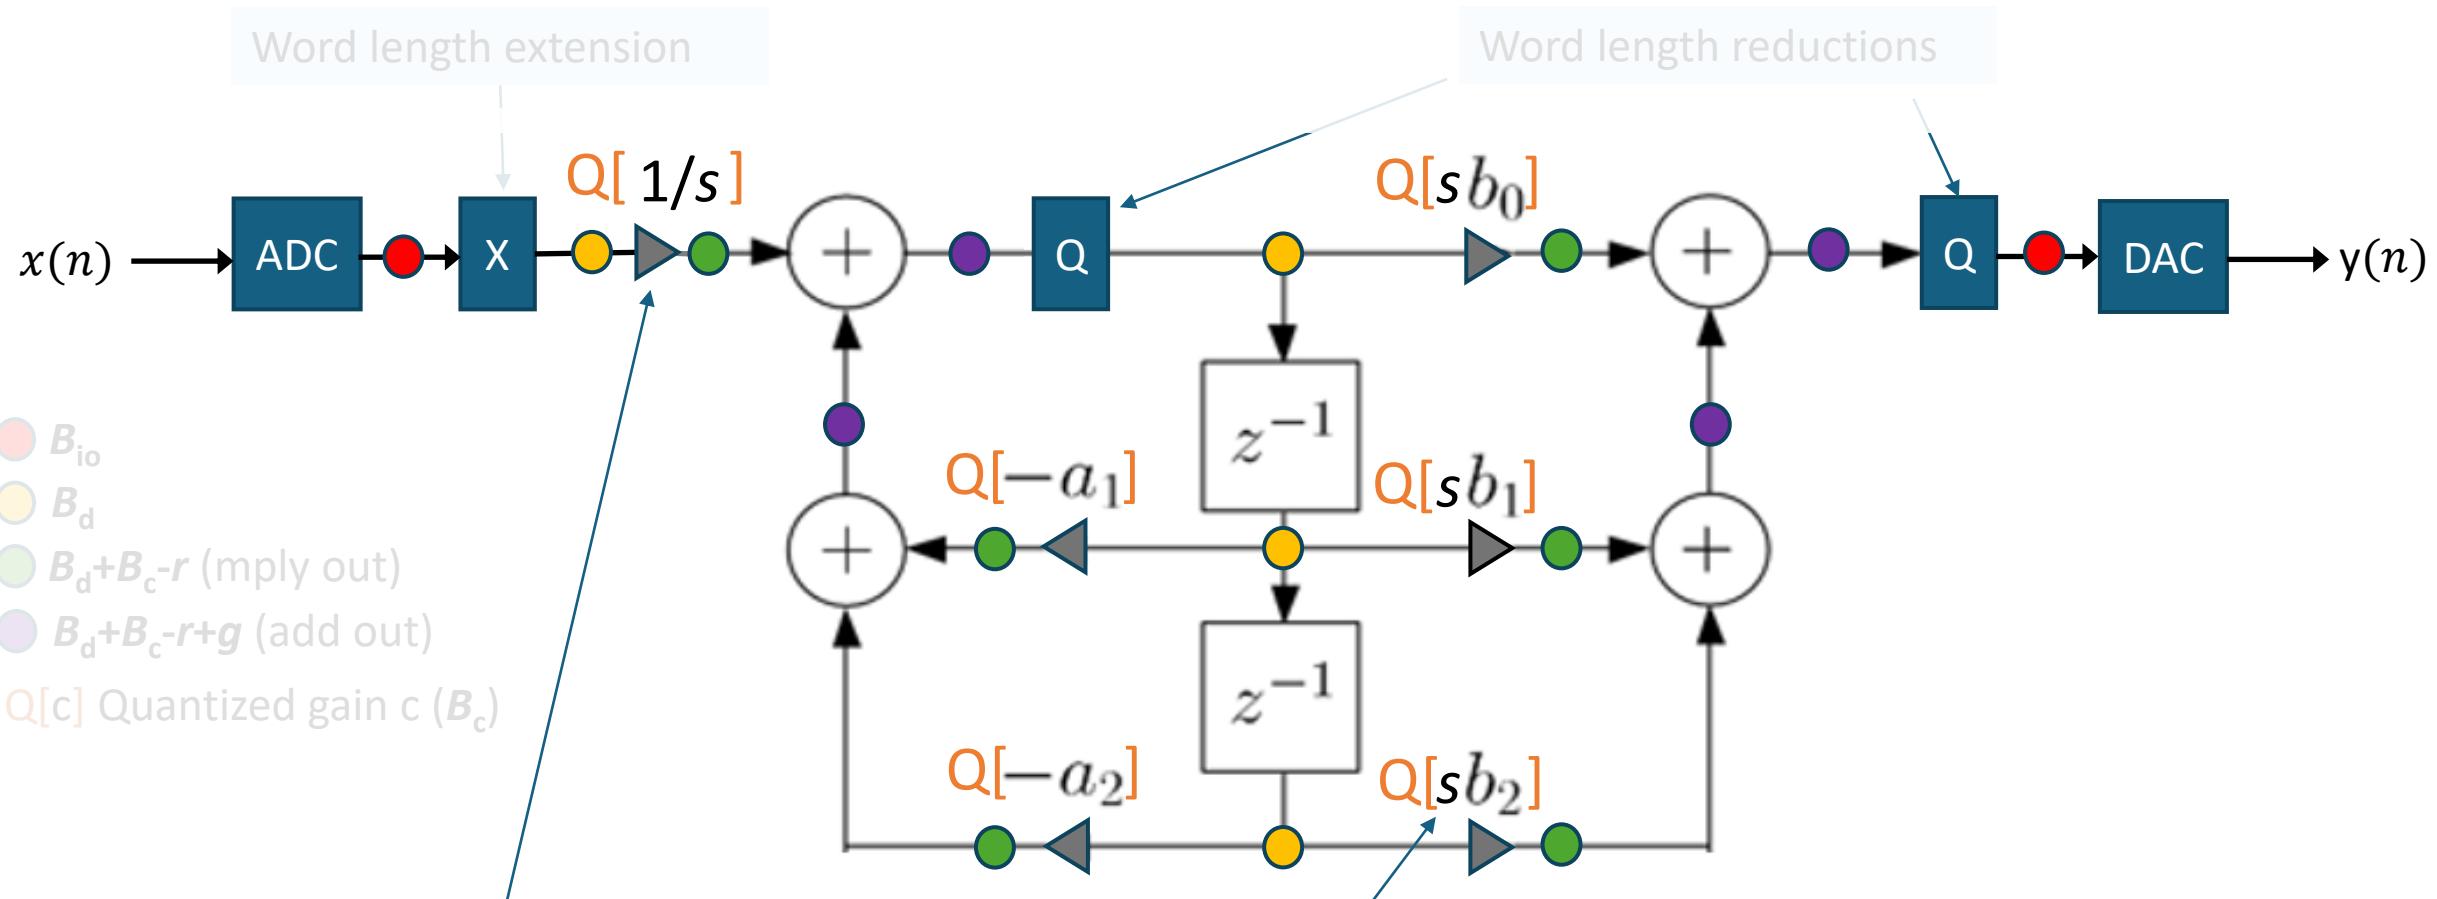


- Assumed DSP platform:
 - Integer arithmetic
 - ● B_{io} -bit I/O
 - ● B_d -bit registers for delays
 - ● B_c -bit memories for coefficients (incl. pre-scaling factor)
 - ● Multiplier output is (B_d+B_c-r) -bit (r is the number of LSB bits discarded from the result by rounding)
 - ● Sums stored into (B_d+B_c-r+g) -bit accumulators (g is the number of guard bits); saturating adder used

Signal flow with DSP platform details



Pre-scaling

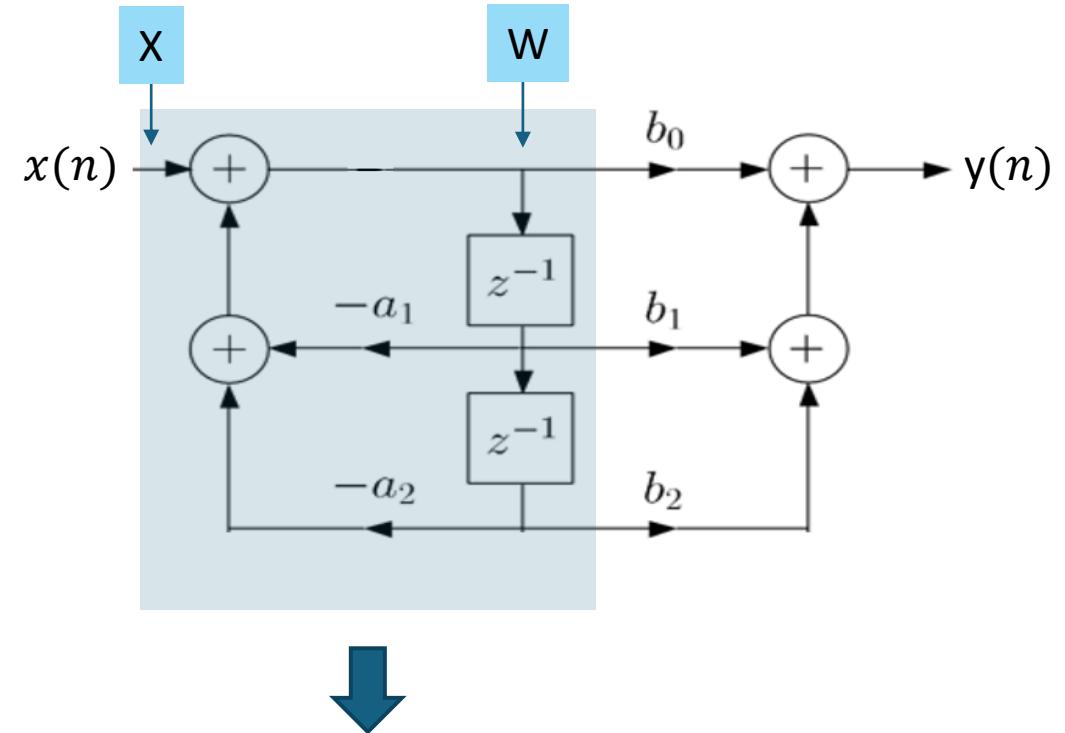


One additional multiplication:

$1/s$ A **pre-scaling** factor (<1) that is used to reduce/avoid overflows in the feedback portion
Compensation for it is done in scaling of feedforward coefficients

3.2. Scaling for overflow control

- $x(n) \in [-1, +1]$
- We reduce the input range so that overflows in the feedback portion are reduced or even avoided completely
- To find the pre-scaling $1/s$, we consider the response at point **W** to the impulse at point **X**
 - Transfer function $H'(z)$
 - A norm of its impulse response $f(n)$ provides the gain s that we compensate for
 - There are various norms that can be applied, computed in time or frequency domain



$$H'(z) = \frac{W(z)}{X(z)} = \frac{1}{1 + a_1 z^{-1} + a_2 z^{-2}}$$

Three common norms for scaling

- L_1 norm applied to impulse response $f(n)$
 - $s = \sum_{n=0}^{\infty} |f(n)|$
 - Overflow avoided, but may be too conservative
- L_2 norm applied to impulse response
 - $s = \sqrt{\sum_{n=0}^{\infty} f^2(n)}$
 - Provides better SNR than previous one (if overflows do not occur)
- L_∞ norm applied to frequency response $F(w)$
 - $s = \max |F(w)|$
 - Overflow with full scale sinusoidal input avoided
- Relationship: $L_2 < L_\infty < L_1$

L2 : Closed form solution for SOS

$$s = \sqrt{\frac{1}{1-a_2^2 - a_1^2(1-a_2)/(1+a_2)}}$$

Matlab:

```
f = impz(1, [1, a1, a2])
L1 = norm(f, 1)
L2 = norm(f, 2)
F = freqz(1, [1, a1, a2])
Linf = max(abs(F))
```

Note:

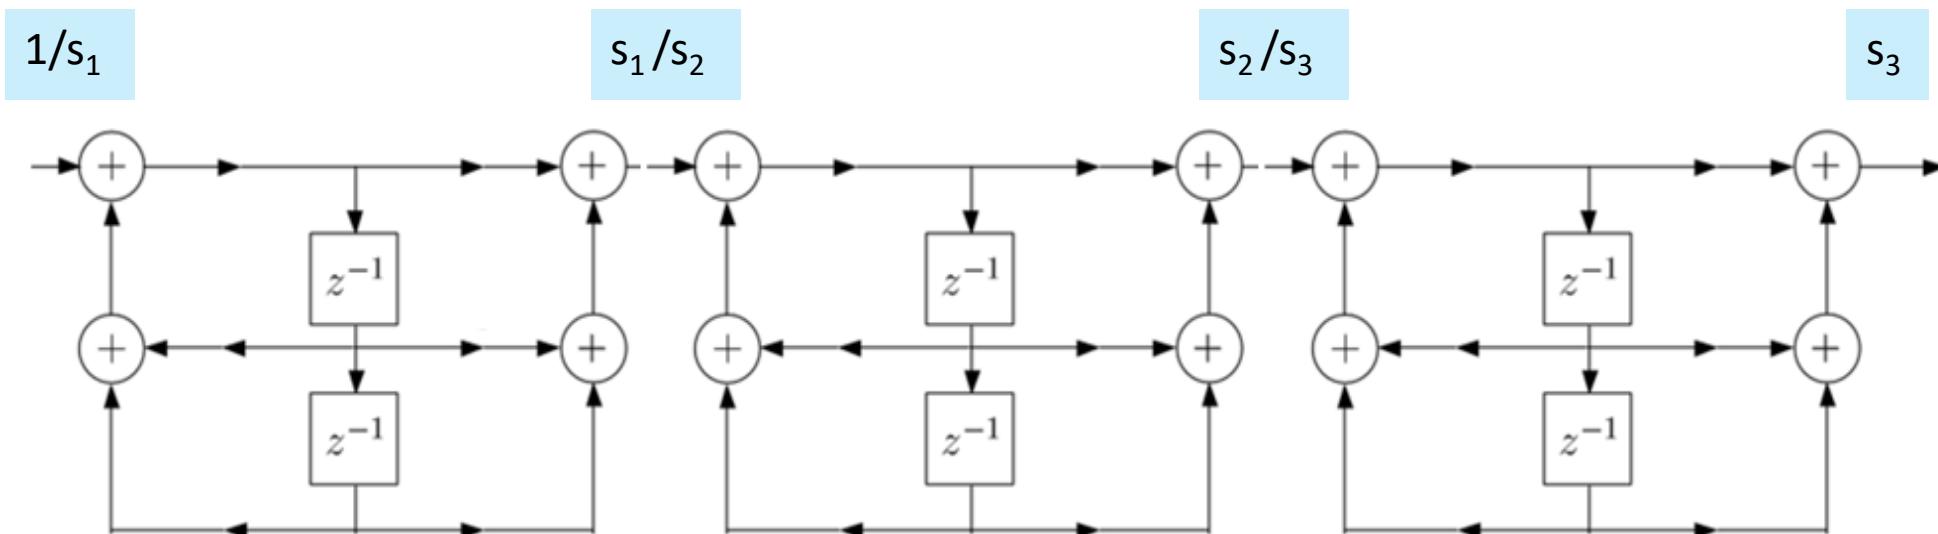
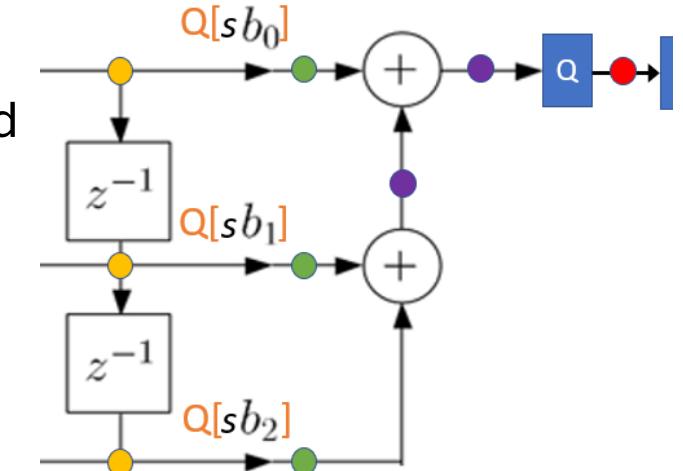
```
L2 = norm(F, 2) / sqrt(numel(F))
```

Notes on scaling

COMPENSATION for $1/s$?

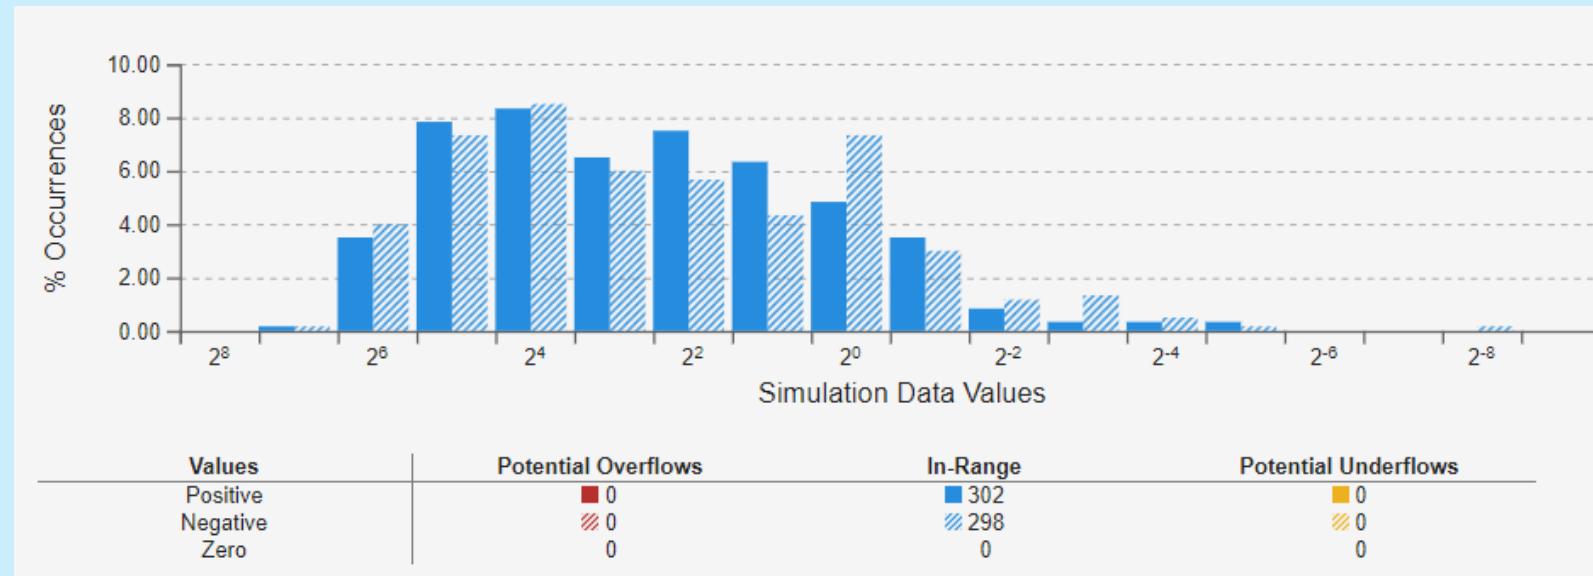
Alternatives

- Scaling of feedforward coefficients b_i by $s \Rightarrow$ filter gain unchanged
- A scaling factor $\neq s$ might be used for them
 - What is the range of the multiplier/accumulator outputs?
- Scaling of b_i might also be neglected
 - Do scaling before rounding to output word length (extra multiplication)
 - In SOS cascades, integrate it to pre-scaling of the following SOS



Tool: NumericTypeScope

- A tool in Fixed Point Toolbox
- Can visualize what kind of values observed for a variable
 - If it has fixed-point type, provides information about overflows and underflows => information on adjusting word and fraction lengths

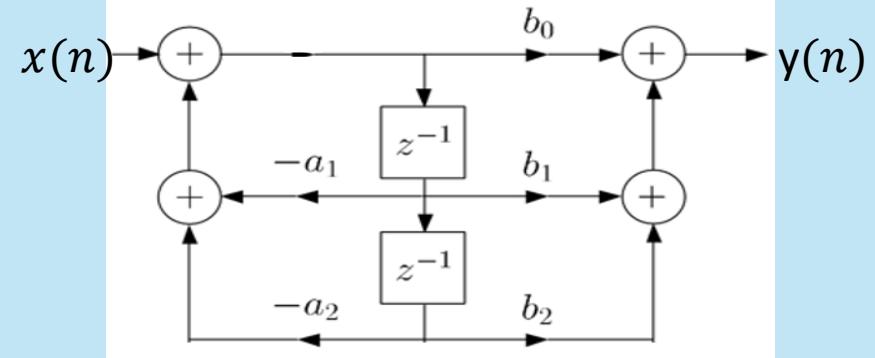


Interpretation: the bars labeled 2^k corresponds to observed magnitudes $2^k \leq |x| < 2^{k+1}$

Input scaling

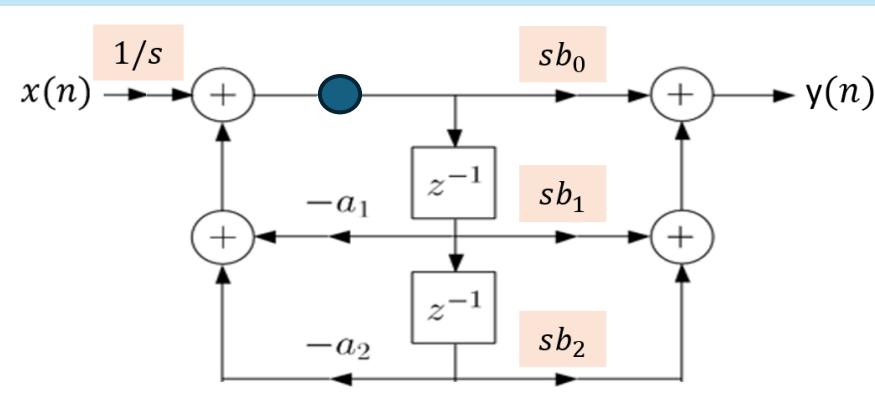
The idea is to prevent overflows in feedback section (e.g., keep results below 1)

```
for k=1:length(x)
    z0 = x(k) - a1 * z1 - a2 * z2;
    y(k) = b0 * z0 + b1 * z1 + b2 * z2;
    z2 = z1; z1 = z0;
end
```



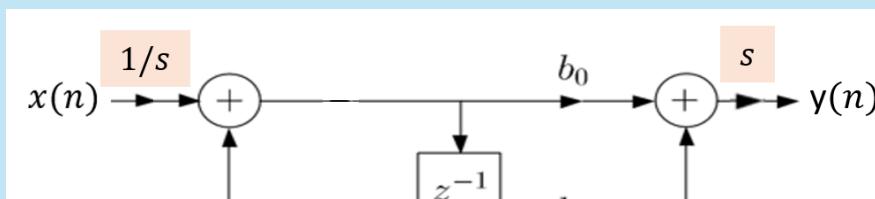
pre-scaling, compensation by scaled feedforward coefficient

```
for k=1:length(x)
    z0 = si * x(k) - a1 * z1 - a2 * z2;
    y(k) = sb0 * z0 + sb1 * z1 + sb2 * z2;
    z2 = z1; z1 = z0;
end
```



pre-scaling, compensation by post-multiplication

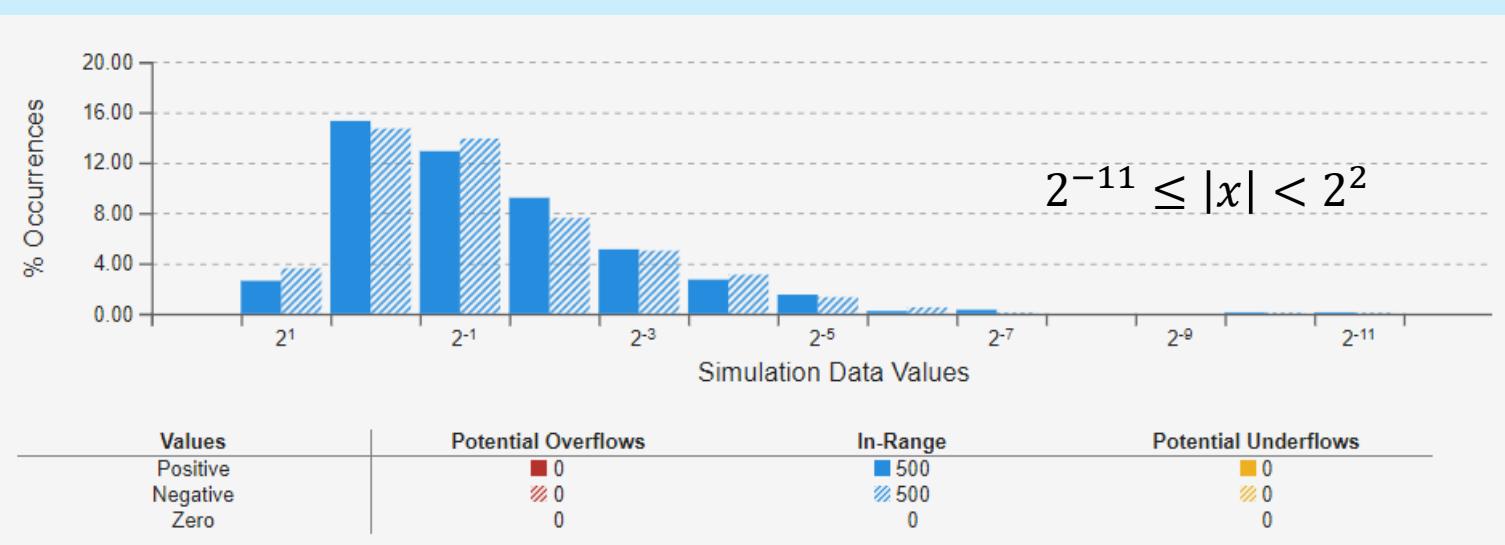
```
for k=1:length(x)
    z0 = si * x(k) - a1 * z1 - a2 * z2;
    siy = b0 * z0 + b1 * z1 + b2 * z2;
    y(k) = s * siy;
    z2 = z1; z1 = z0;
end
```



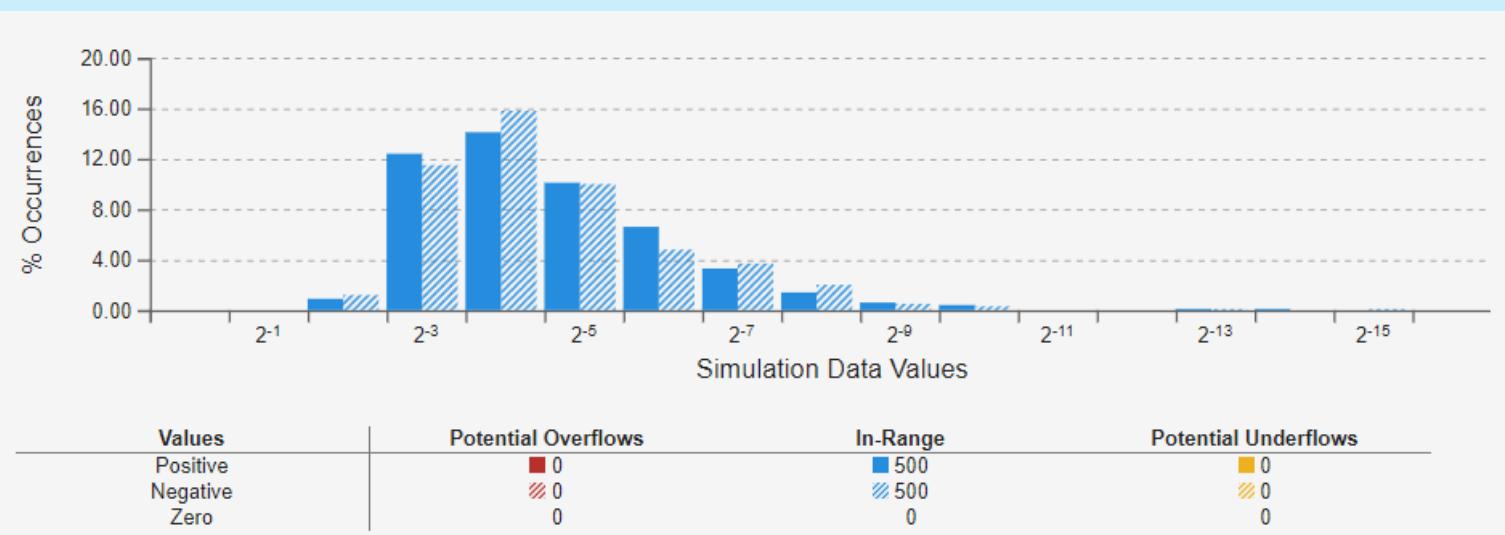
Example

- scopeDemo mlx
- Delay register variable ranges
- If one needs to keep delay register values below 1, probably latter solution would work

No scaling



Prescaling with $s = 10$



Range

$2^{-15} \leq |x| < 2^{-1}$

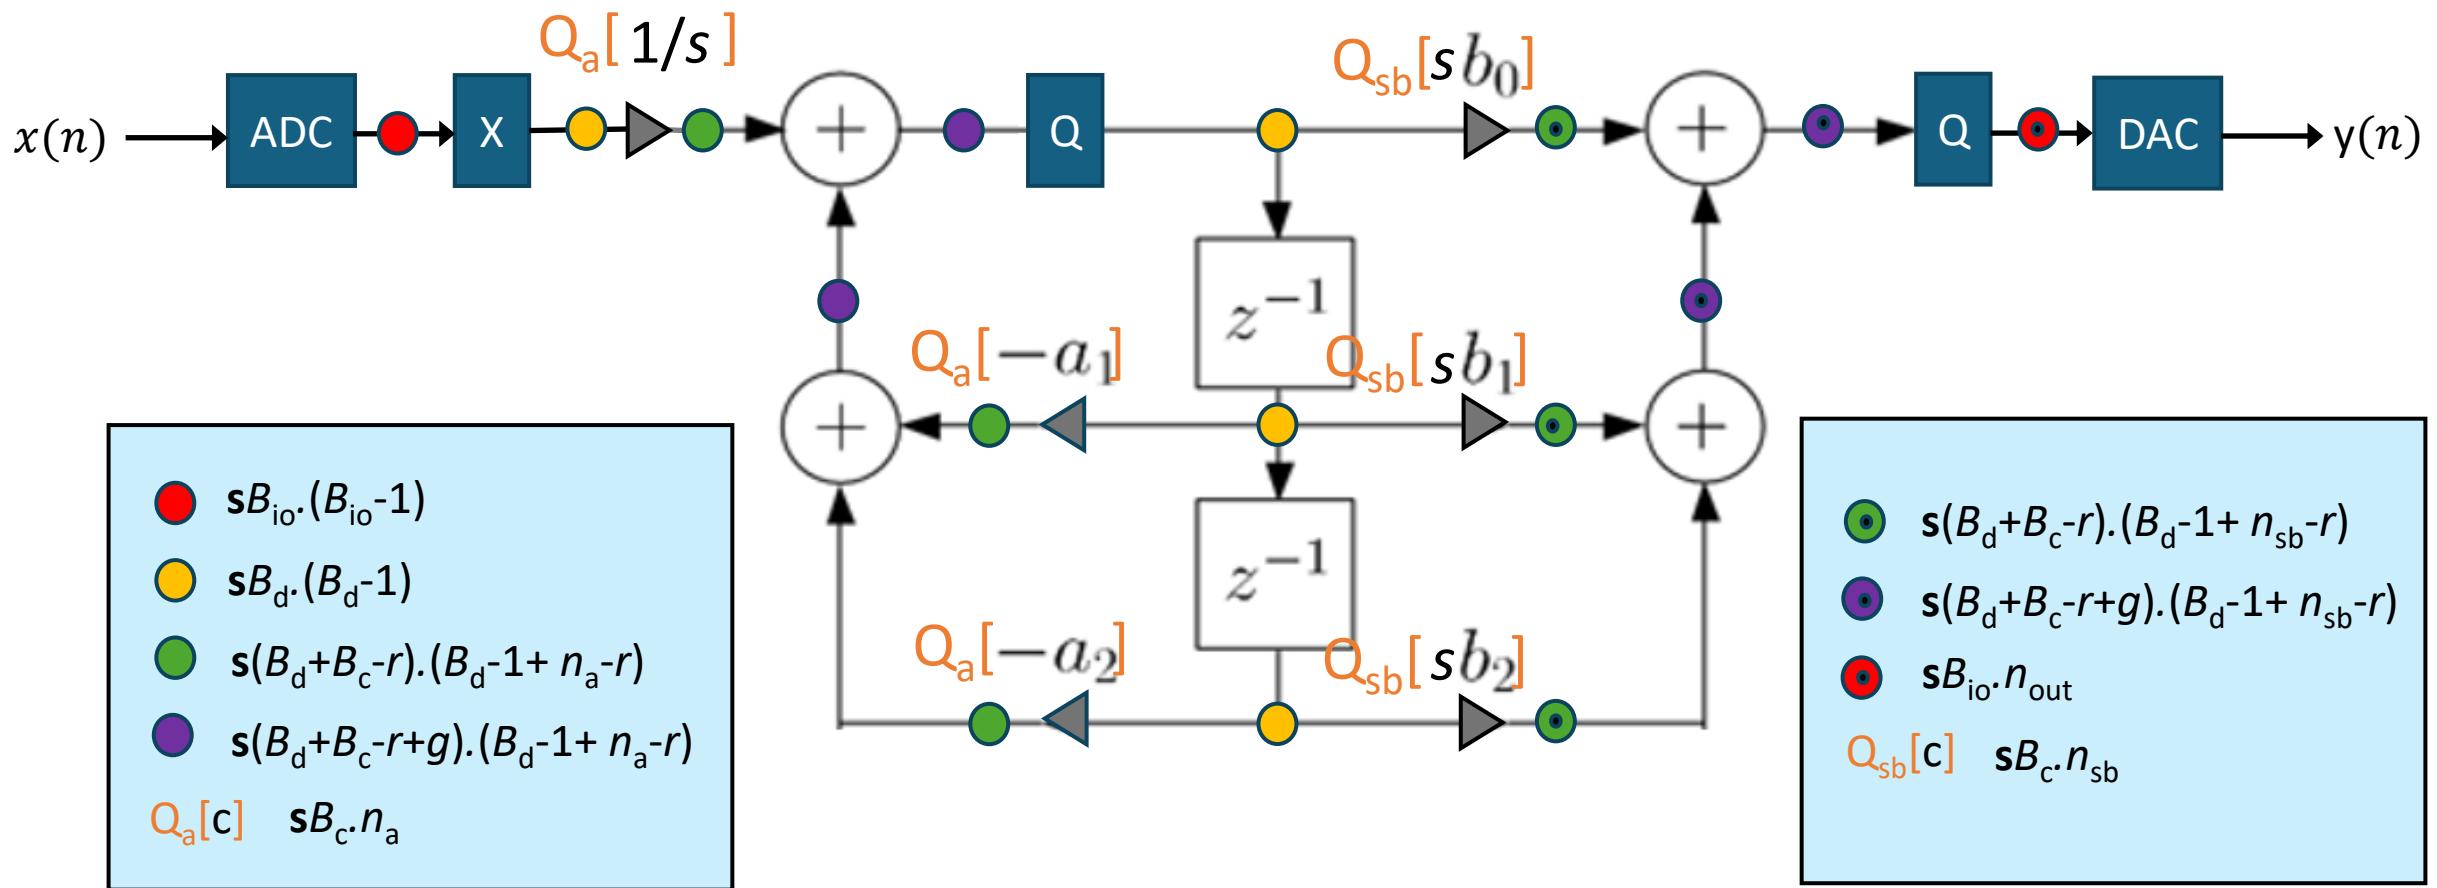
3.3. Data path: fixed-point formats

- In previous, only **word lengths** of DSP platform were given
- We must make decisions related to the **fraction lengths**
 - In some cases, we can infer them
 - Some decisions require simulation studies
- Fraction lengths are defined for
 - 1) Input: assuming range (-1,+1) so the fraction length is $B_{io}-1$
 - 2) Output: set by user, n_{out} depends on the observed gain of the filter
 - 3) Delay-line: pre-scaling => range (-1,+1) assumed so the fraction length is B_d-1
 - 4) Coefficients: separate choice of fraction length (n_a , n_{sb}) for feedback/-forward portions
 - Guided by the range of coefficients $\{1/s, a_1, a_2\}$ and $\{sb_0, sb_1, sb_2\}$
 - 5) Multiplier output: choice of fraction length is guided by inputs and amount of LSB reduction r
 - 6) Adder output: follows from the multiplier output format
 - one or two guard bits g may be added, but they do not affect the fraction length



●	B_{io}
●	B_d
●	B_d+B_c-r (mply out)
●	B_d+B_c-r+g (add out)
Q[c]	Quantized gain c (B_c)

Signal flow with fixed-point sp.n formats



Summary of noise sources

- Multiplications 

- Source of round-off errors
- Only if $r > 0$, that is, if we drop LSBs by truncation (or some other rounding)

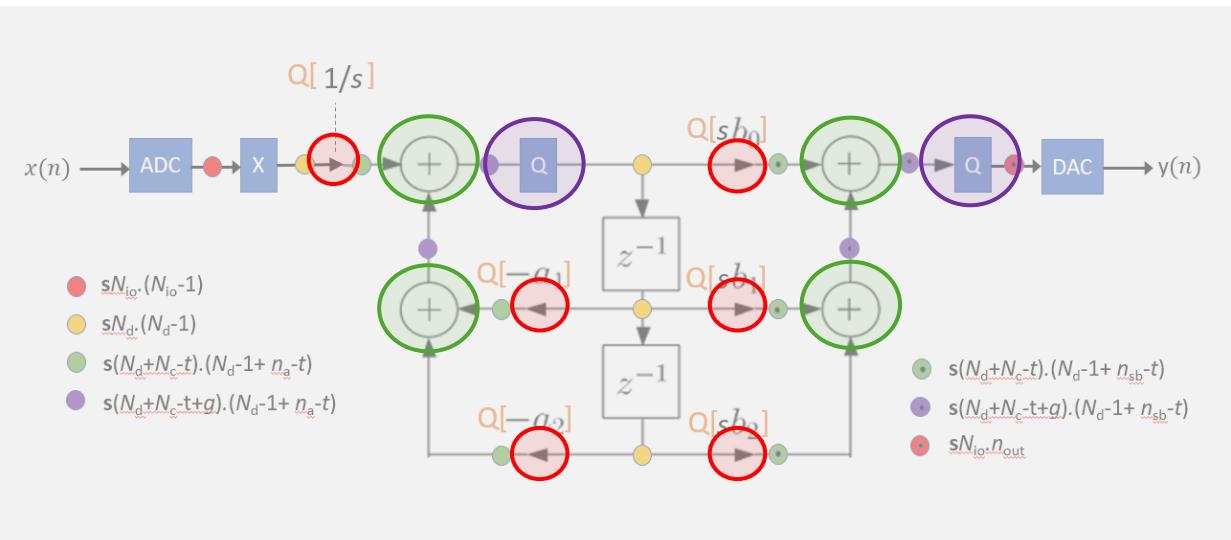
- Additions 

- Potential source of overflow errors

- Conversions 

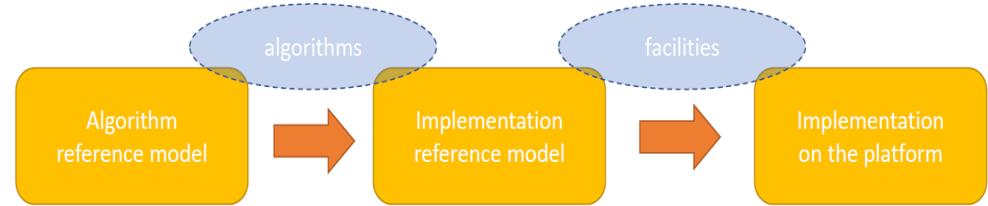
- Q blocks before the delay line and DAC
- round-off errors, overflow errors possible

- In addition, coefficient quantizations have effect on DSP noise (change response from optimal)



3.4. Simulator for experiments

- To experiment with the implementation outlined in the preceding, some Matlab code were written, utilizing Fixed Point Toolbox
 - The code, **SOSSimulator.m**, is provided in zip package
 - Example of an **implementation reference model**



```
% SOSSimulator    Simulate a second-order section
%
% << Syntax >>
%   h = SOSSimulator(b,a,sp,[Bi,Bc,Bd],r,g,no,mround,qround)
%
% << Arguments >>
%   b           Numerator coefficients : [b(0),b(1),b(2)]
%   a           Denominator coefficients : [a(0),a(1),a(2)]
%   sp          Scaling factor : computation method ('L1t', 'L2t', 'Lif') or a numeric value
%               - note: coefficients "b" are multiplied by this factor, input by 1/sp
%   Bi,Bc,Bd  Word length used for I/O, coefficients and delay memories
%               - multiplier outputs have word length Bc+Bd
%               - adder outputs have word length Bc+Bd+g
%   r           Number of bits discarded at multiplier output
%               - if r > 0, there will be round-off errors in multiplications
%   g           Number of accumulator guard bits (0-2)
%               - if necessary, can be used to prevent overflows in additions
%   no          Number of fraction bits used for output
%   mround     Multiplier rounding method : 'floor', 'nearest', 'round', 'zero' or 'convergent'
%               - see fimath RoundingMethod for meaning
%               - note: 'floor' is same as truncation, which would be easiest to implement in HW
%   qround     Rounding method for word length reduction (Q) blocks
%   h           Output: structure of access function handles
```

SOSsimulator.m's procedure for setting up a fixed-point implementation for SOS filter

1. Quantize the feedback coefficients

- * **Design parameter:** word length used for coefficients (p_c)
- * Determine the fraction length under the assumption that all feedback coefficients must use the same format
- * Quantize coefficients a_1 and a_2

2. Set the scaling factor

- * The transfer function analyzed is $H'(z) = \frac{1}{1+a_1z^{-1}+a_2z^{-2}}$.
- * **Design parameter:** Select one of the norms of $H'(z)$ or provide some other value
- * Quantize pre-scaling factor 1/s using the format used for feedback coefficients

3. Scale and quantize the feedforward coefficients

- * **Design parameter:** word length used for coefficients (p_c)
- * To compensate for pre-scaling, multiply feedforward coefficients by inverse of the quantized pre-scaling factor
- * Determine the fraction length under the assumption that all scaled feedforward coefficients must use the same format
- * Quantize coefficients sb_0 , sb_1 and sb_2

4. Determine fixed-point formats on the data paths

- * **Design parameters:** p_{io} , p_d , p_c , t , g , and n_o
- * Input format is $sp_{io}.(p_{io}-1)$, where p_{io} is the word length for input/output
- * Delay memory format is $sp_d.(p_d-1)$, where p_d is the word length for internal memory
- * Multiplier output formats: $s(p_c + p_d - t).(n + p_d - 1 - t)$, where n depends on the portion (feedback/-forward) of the filter
- * Accumulator formats: $s(p_c + p_d - t + g).(n + p_d - 1 - t)$, where g is the number of guard bits used
- * Output format is $sp_{io}.n_o$, where n_o is the fraction length chosen by the user

Simulate using test signals and repeat with other design parameters, if not satisfactory

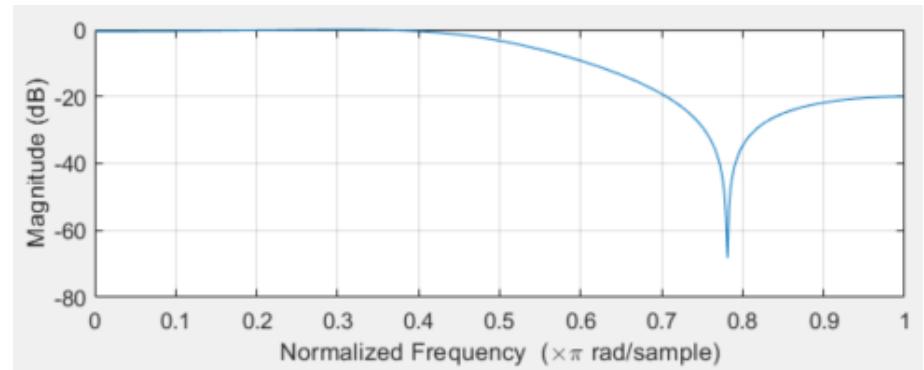
- * Simulation interests: Overflows and underflows at acceptable level? SNR acceptable? Resource use optimal?
- * Compare to floating-point!

3.5. Experiments*

In the experiments, a low-pass SOS filter is constructed with the command

```
>> [b,a] = ellip(2,0.5,20,0.4);
```

Frequency response of the filter on the right. Passband cutoff is at $0.4 \times F_s / 2$, stopband cutoff at $0.7 \times F_s/2$, and stopband attenuation 20 dB.



Three experiments:

Experiment 1 – truncation versus rounding

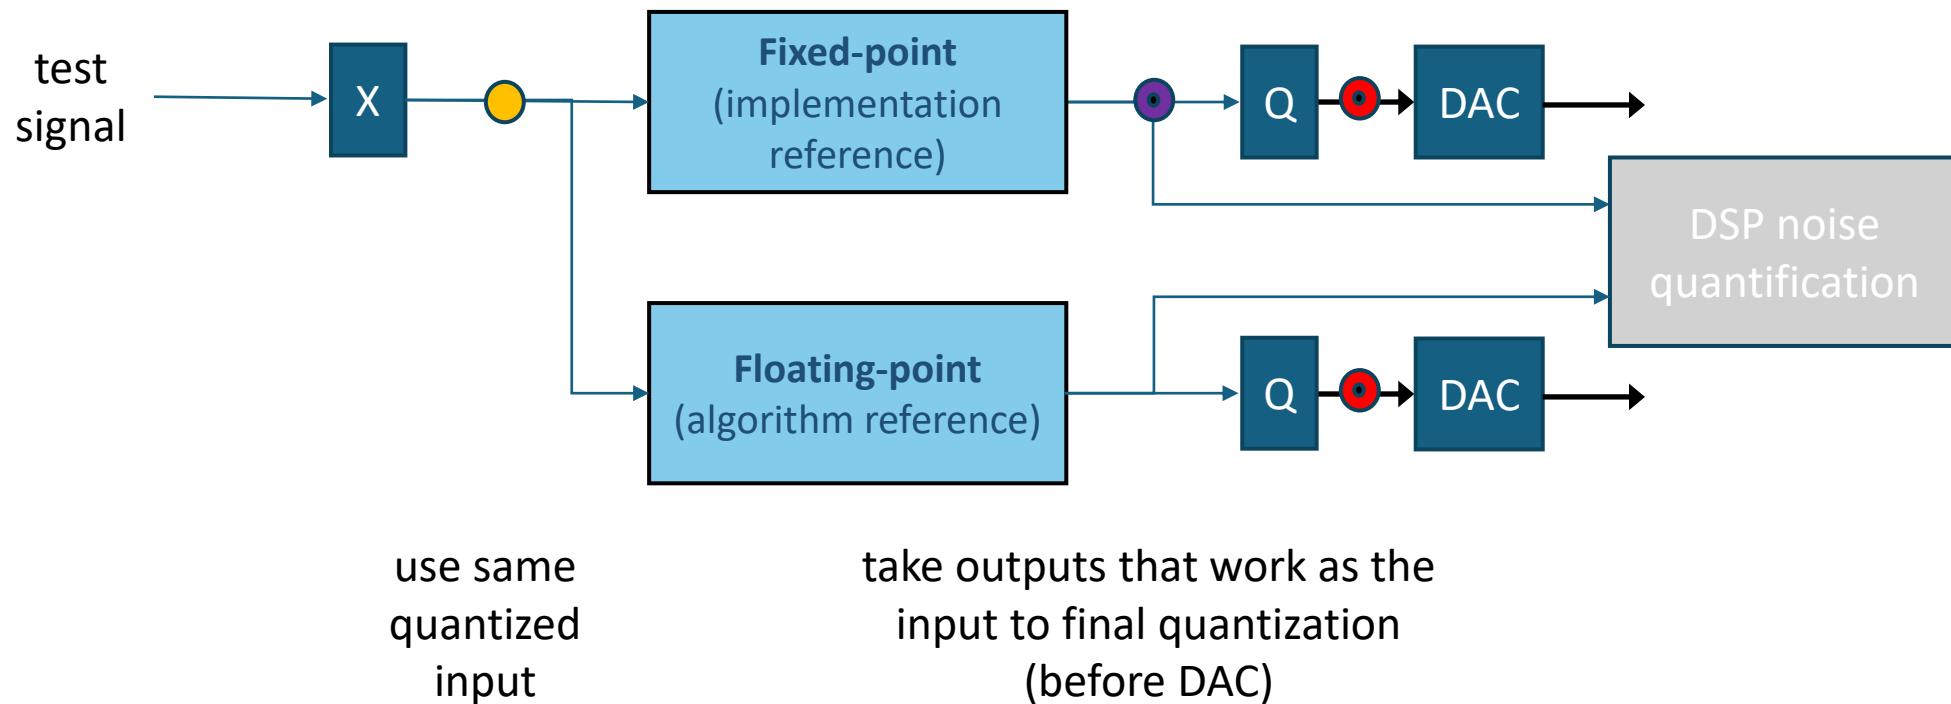
Experiment 2 – required precision of multiplications

Experiment 3 – error correlations with periodic signals

Fixed-point constructs are compared to the full-precision floating-point filter, which can be simulated with Matlab's **filter** command.

Evaluation setup (1)

Following setup is used in order to make output comparison fair



Evaluation setup (2)

- **Sine wave test signal**

- Varying frequency within the range (0, $F_s/2$)

$$x(n) = \sin(2\pi f n + r)$$

where $f \in (0, 1/2)$ and r is a small random phase shift

- **ADC noise at output**

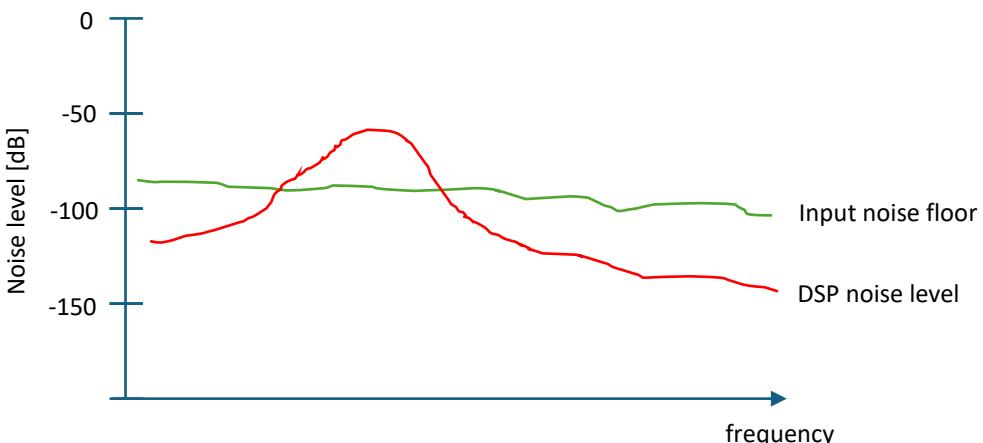
- Evaluated in order to compare the DSP noise to ADC noise
- Model-based formula used

- Plotting ratio $\frac{\text{noise power}}{\text{signal power}}$ [dB]

$$\sigma_{OA}^2 = \sigma_A^2 \sum_{k=0}^{\infty} h^2(k) \quad \sigma_A^2 = \frac{2^{-2B}}{3}$$

where $h(k)$ is impulse response of the filter and B is the ADC word length.

- Recall earlier noise floor figure:

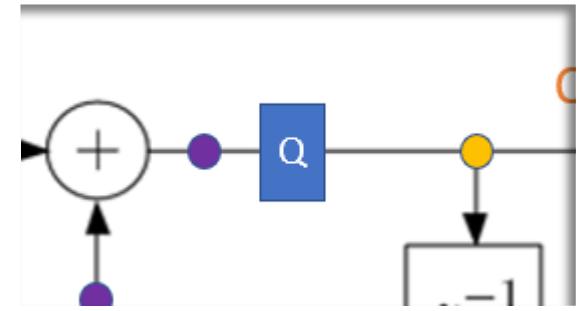
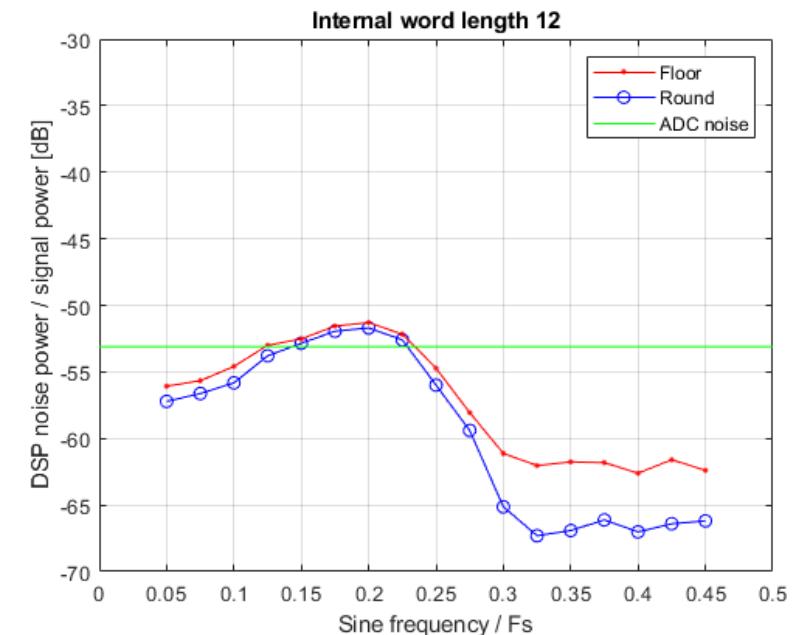
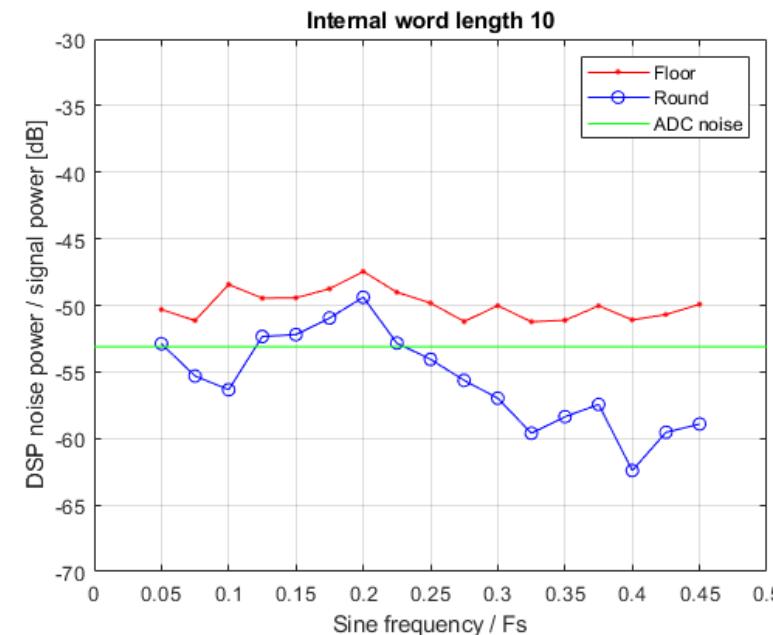
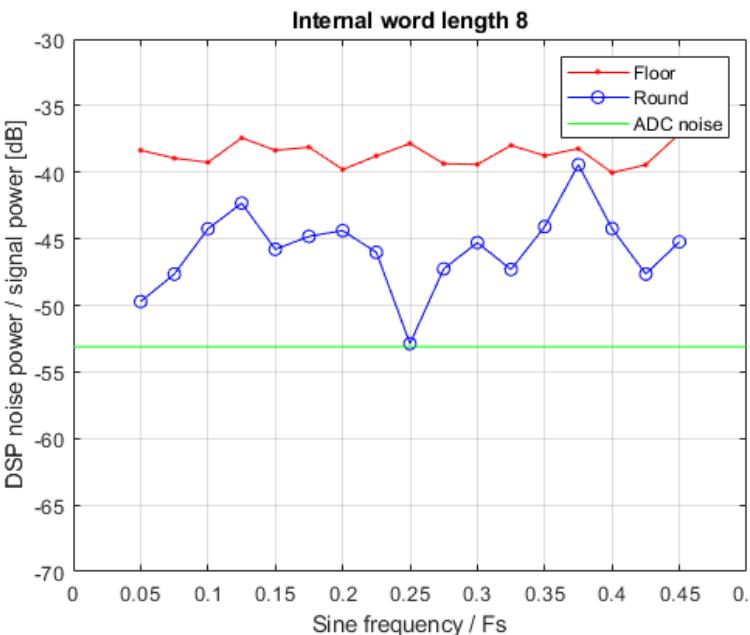


*Code for following experiments, ***SOSexperiments.m***, is in zip package

Experiment 1 – truncation study

- Quantization by truncation (= floor) is biased. Does it show up?
- Small experiment, where we compare truncation to rounding
- 8-bit ADC, 8-bit coefficients
- **delay register word length (N_d)** varied from 8 to 12 bits
- Result:

$$N_d = 8$$

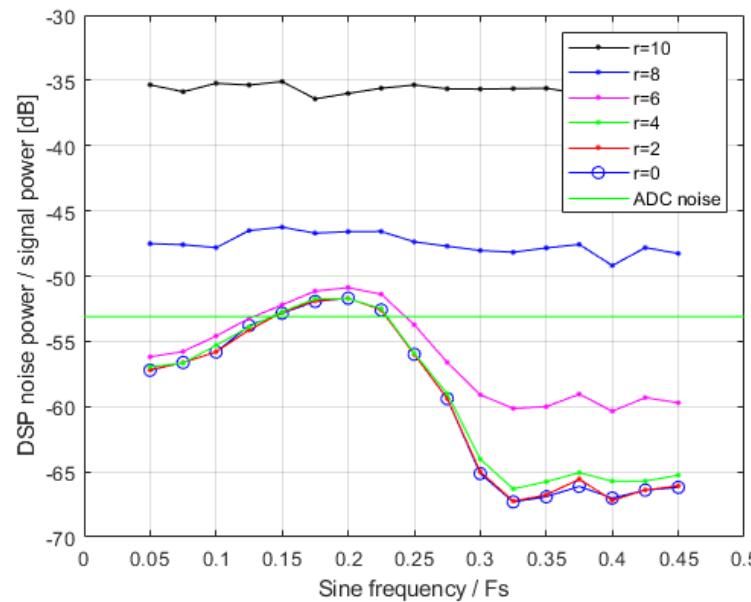


Truncation/rounding point

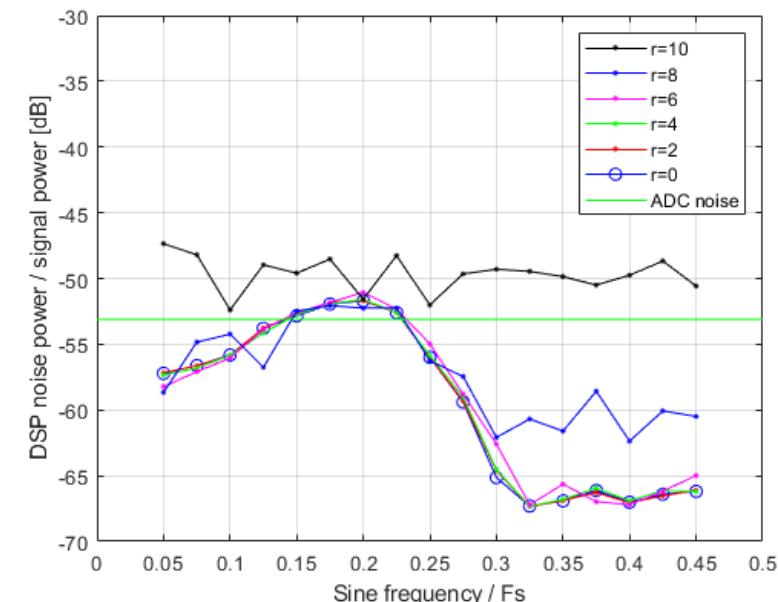
Experiment 2 – precision of multiplications

- How much can we discard LSB bits in multiplication?
- Experiment where we increase the parameter r , which controls discarding
- Delay register word length set to 12 bits, Q operation uses rounding
- Multiplier output word length $(20 - r)$, testing for $r = 0, 2, 4, 6, 8, 10$ (floor / round)
- Result:

Floor:



Round:



More bits can be discarded in the case of rounding. Compare results with $r = 8$

Experiment 3 – error correlations

- DSP error spectra of periodic and aperiodic signals compared

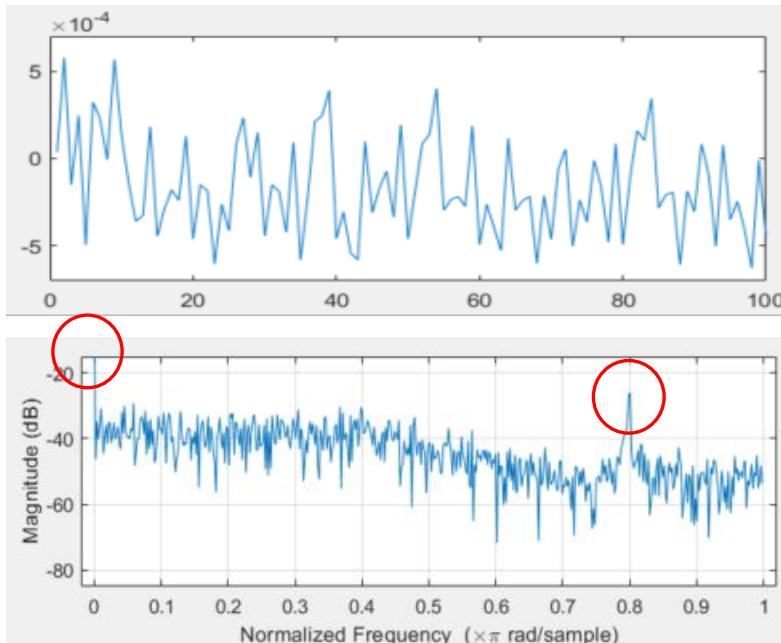
- The test signal

$$x(n) = \sin(2\pi fn + r)$$

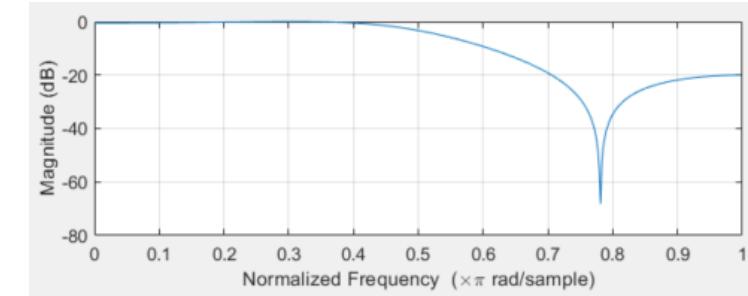
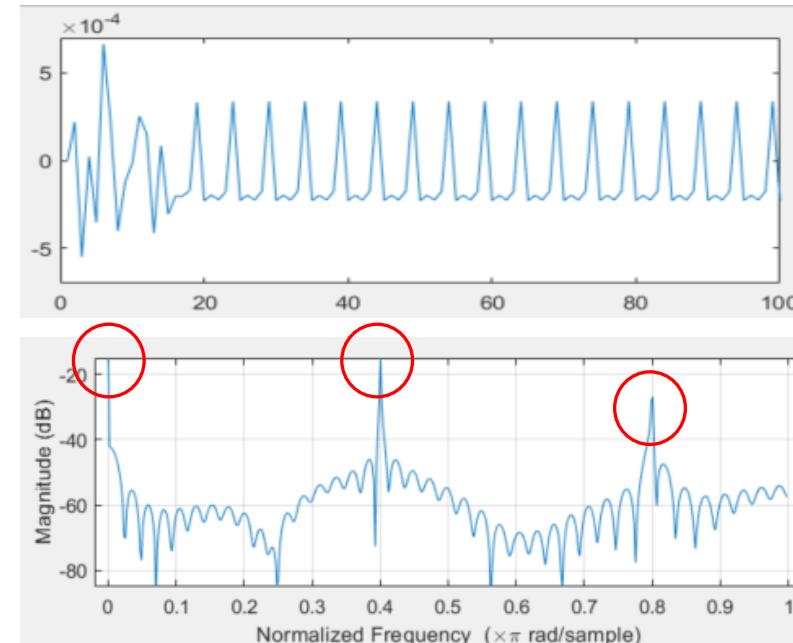
where $f \in (0, 1/2)$ and r is a phase shift

- Comparison for sine frequency $f = 0.4$:

With **random** phase shift r



With **zero** phase shift $r (r=0)$



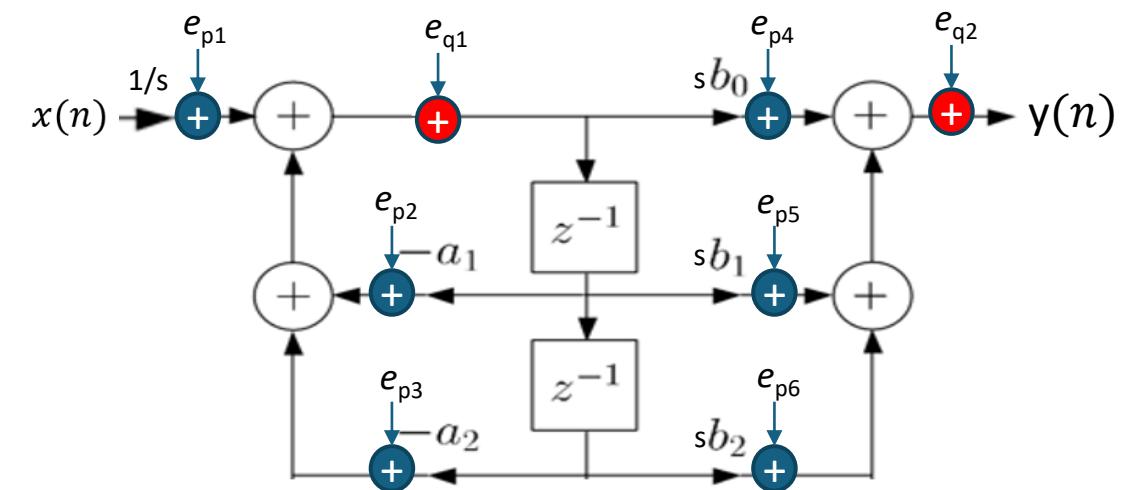
Response of the filter

Error correlations cause peaks to the error spectra

Error at zero frequency : dc bias in output

3.6. Model-based analysis*

- The alternative for simulation-based approach was the analytic approach
- In analysis, noise sources have certain assumed statistics (spectrally white, independent)
- Effect at the filter output computed
 - the impulse response from the noise source to the output must be determined
 - Uncorrelated => effects of noise sources summed up
- Note 1: moving an added input across a sum has no effect
 - Direct effect of e_{p4} , e_{p5} , e_{p6} , e_{q2} at output
- Note 2: many noise sources have the same transfer function
 - e_{p1} , e_{p2} , e_{p3} , e_{q1} have the same path through the filter
- See Ifeachor & Jervis (2002) for in-depth discussion



= noise due to word length reduction

= noise due to product quantization ($r > 0$)

Summary

- A/D conversion
 - Tuning – Finding balance between quantization and clipping noises
 - Noise floor – Setting target for quality of signal processing
- IIR filter coefficient quantization
 - Filter stability must be addressed
 - Need to consider structure of computation
- IIR filter processing
 - Multiple fixed-point formats involved
 - Scaling of input data helps to prevent overflows in feedback sections
 - NumericTypeScope in Matlab FPT provides a tool for analyzing data collected during floating-point or fixed-point simulation